

Data Manipulation with pandas

Transforming Data

Let's master the pandas basics. Learn how to inspect DataFrames and perform fundamental manipulations, including sorting rows, subsetting, and adding new columns.

[Link for reference](#)

Inspecting a DataFrame

```
In [ ]: # Import pandas using the alias pd
import pandas as pd

#pathway of the file
homelessness = pd.read_csv('C:\\Users\\yeiso\\OneDrive - Douglas College\\0. DOUGLAS COLLEGE\\3. Fund Machine Learning\\0. Python
```

```
In [ ]: # Print the head of the homelessness data
print(homelessness.head())
```

	Unnamed: 0	region	state	individuals	family_members	\
0	0	East South Central	Alabama	2570.0	864.0	
1	1	Pacific	Alaska	1434.0	582.0	
2	2	Mountain	Arizona	7259.0	2606.0	
3	3	West South Central	Arkansas	2280.0	432.0	
4	4	Pacific	California	109008.0	20964.0	

	state_pop
0	4887681
1	735139
2	7158024
3	3009733
4	39461588

```
In [ ]: # Print information about homelessness
print(homelessness.info())
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 51 entries, 0 to 50
Data columns (total 6 columns):
#   Column          Non-Null Count  Dtype
---  -
0   Unnamed: 0      51 non-null    int64
1   region          51 non-null    object
2   state           51 non-null    object
3   individuals     51 non-null    float64
4   family_members  51 non-null    float64
5   state_pop       51 non-null    int64
dtypes: float64(2), int64(2), object(2)
memory usage: 2.5+ KB
None

```

```

In [ ]: # Print the shape of homelessness
print(homelessness.shape)

```

```

(51, 6)

```

```

In [ ]: # Print a description of homelessness
print(homelessness.describe())

```

	Unnamed: 0	individuals	family_members	state_pop
count	51.000000	51.000000	51.000000	5.100000e+01
mean	25.000000	7225.784314	3504.882353	6.405637e+06
std	14.866069	15991.025083	7805.411811	7.327258e+06
min	0.000000	434.000000	75.000000	5.776010e+05
25%	12.500000	1446.500000	592.000000	1.777414e+06
50%	25.000000	3082.000000	1482.000000	4.461153e+06
75%	37.500000	6781.500000	3196.000000	7.340946e+06
max	50.000000	109008.000000	52070.000000	3.946159e+07

Parts of a DataFrame

```

In [ ]: # Print the values of homelessness
print(homelessness.values)

```

[[0 'East South Central' 'Alabama' 2570.0 864.0 4887681]
[1 'Pacific' 'Alaska' 1434.0 582.0 735139]
[2 'Mountain' 'Arizona' 7259.0 2606.0 7158024]
[3 'West South Central' 'Arkansas' 2280.0 432.0 3009733]
[4 'Pacific' 'California' 109008.0 20964.0 39461588]
[5 'Mountain' 'Colorado' 7607.0 3250.0 5691287]
[6 'New England' 'Connecticut' 2280.0 1696.0 3571520]
[7 'South Atlantic' 'Delaware' 708.0 374.0 965479]
[8 'South Atlantic' 'District of Columbia' 3770.0 3134.0 701547]
[9 'South Atlantic' 'Florida' 21443.0 9587.0 21244317]
[10 'South Atlantic' 'Georgia' 6943.0 2556.0 10511131]
[11 'Pacific' 'Hawaii' 4131.0 2399.0 1420593]
[12 'Mountain' 'Idaho' 1297.0 715.0 1750536]
[13 'East North Central' 'Illinois' 6752.0 3891.0 12723071]
[14 'East North Central' 'Indiana' 3776.0 1482.0 6695497]
[15 'West North Central' 'Iowa' 1711.0 1038.0 3148618]
[16 'West North Central' 'Kansas' 1443.0 773.0 2911359]
[17 'East South Central' 'Kentucky' 2735.0 953.0 4461153]
[18 'West South Central' 'Louisiana' 2540.0 519.0 4659690]
[19 'New England' 'Maine' 1450.0 1066.0 1339057]
[20 'South Atlantic' 'Maryland' 4914.0 2230.0 6035802]
[21 'New England' 'Massachusetts' 6811.0 13257.0 6882635]
[22 'East North Central' 'Michigan' 5209.0 3142.0 9984072]
[23 'West North Central' 'Minnesota' 3993.0 3250.0 5606249]
[24 'East South Central' 'Mississippi' 1024.0 328.0 2981020]
[25 'West North Central' 'Missouri' 3776.0 2107.0 6121623]
[26 'Mountain' 'Montana' 983.0 422.0 1060665]
[27 'West North Central' 'Nebraska' 1745.0 676.0 1925614]
[28 'Mountain' 'Nevada' 7058.0 486.0 3027341]
[29 'New England' 'New Hampshire' 835.0 615.0 1353465]
[30 'Mid-Atlantic' 'New Jersey' 6048.0 3350.0 8886025]
[31 'Mountain' 'New Mexico' 1949.0 602.0 2092741]
[32 'Mid-Atlantic' 'New York' 39827.0 52070.0 19530351]
[33 'South Atlantic' 'North Carolina' 6451.0 2817.0 10381615]
[34 'West North Central' 'North Dakota' 467.0 75.0 758080]
[35 'East North Central' 'Ohio' 6929.0 3320.0 11676341]
[36 'West South Central' 'Oklahoma' 2823.0 1048.0 3940235]
[37 'Pacific' 'Oregon' 11139.0 3337.0 4181886]
[38 'Mid-Atlantic' 'Pennsylvania' 8163.0 5349.0 12800922]
[39 'New England' 'Rhode Island' 747.0 354.0 1058287]
[40 'South Atlantic' 'South Carolina' 3082.0 851.0 5084156]
[41 'West North Central' 'South Dakota' 836.0 323.0 878698]
[42 'East South Central' 'Tennessee' 6139.0 1744.0 6771631]
[43 'West South Central' 'Texas' 19199.0 6111.0 28628666]
[44 'Mountain' 'Utah' 1904.0 972.0 3153550]
[45 'New England' 'Vermont' 780.0 511.0 624358]
[46 'South Atlantic' 'Virginia' 3928.0 2047.0 8501286]
[47 'Pacific' 'Washington' 16424.0 5880.0 7523869]
[48 'South Atlantic' 'West Virginia' 1021.0 222.0 1804291]]

```
[49 'East North Central' 'Wisconsin' 2740.0 2167.0 5807406]  
[50 'Mountain' 'Wyoming' 434.0 205.0 577601]]
```

```
In [ ]: # Print the column index of homelessness  
print(homelessness.columns)
```

```
Index(['Unnamed: 0', 'region', 'state', 'individuals', 'family_members',  
      'state_pop'],  
      dtype='object')
```

```
In [ ]: # Print the row index of homelessness  
print(homelessness.index)
```

```
RangeIndex(start=0, stop=51, step=1)
```

Sorting rows

```
In [ ]: # Sort homelessness by individuals  
homelessness_ind = homelessness.sort_values(["individuals"])  
  
# Print the top few rows  
print(homelessness_ind.head())
```

	Unnamed: 0	region	state	individuals	family_members	\
50	50	Mountain	Wyoming	434.0	205.0	
34	34	West North Central	North Dakota	467.0	75.0	
7	7	South Atlantic	Delaware	708.0	374.0	
39	39	New England	Rhode Island	747.0	354.0	
45	45	New England	Vermont	780.0	511.0	

	state_pop
50	577601
34	758080
7	965479
39	1058287
45	624358

```
In [ ]: # Sort homelessness by descending family members  
homelessness_fam = homelessness.sort_values(["family_members"], ascending=[False])  
  
# Print the top few rows  
print(homelessness_fam.head())
```

	Unnamed: 0	region	state	individuals	\
32	32	Mid-Atlantic	New York	39827.0	
4	4	Pacific	California	109008.0	
21	21	New England	Massachusetts	6811.0	
9	9	South Atlantic	Florida	21443.0	
43	43	West South Central	Texas	19199.0	

	family_members	state_pop
32	52070.0	19530351
4	20964.0	39461588
21	13257.0	6882635
9	9587.0	21244317
43	6111.0	28628666

```
In [ ]: # Sort homelessness by region, then descending family members
homelessness_reg_fam = homelessness.sort_values(["region", "family_members"], ascending=[True, False])

# Print the top few rows
print(homelessness_reg_fam.head())
```

	Unnamed: 0	region	state	individuals	family_members	\
13	13	East North Central	Illinois	6752.0	3891.0	
35	35	East North Central	Ohio	6929.0	3320.0	
22	22	East North Central	Michigan	5209.0	3142.0	
49	49	East North Central	Wisconsin	2740.0	2167.0	
14	14	East North Central	Indiana	3776.0	1482.0	

	state_pop
13	12723071
35	11676341
22	9984072
49	5807406
14	6695497

Subsetting columns

```
In [ ]: # Select the individuals column
individuals = homelessness["individuals"]

# Print the head of the result
print(individuals.head())
```

0	2570.0
1	1434.0
2	7259.0
3	2280.0
4	109008.0

Name: individuals, dtype: float64

```
In [ ]: # Select the state and family_members columns
state_fam = homelessness[["state", "family_members"]]

# Print the head of the result
print(state_fam.head())
```

	state	family_members
0	Alabama	864.0
1	Alaska	582.0
2	Arizona	2606.0
3	Arkansas	432.0
4	California	20964.0

```
In [ ]: # Select only the individuals and state columns, in that order
ind_state = homelessness[["individuals", "state"]]

# Print the head of the result
print(ind_state.head())
```

	individuals	state
0	2570.0	Alabama
1	1434.0	Alaska
2	7259.0	Arizona
3	2280.0	Arkansas
4	109008.0	California

Subsetting rows

```
In [ ]: # Filter for rows where individuals is greater than 10000
ind_gt_10k = homelessness[homelessness["individuals"] > 10000]

# See the result
print(ind_gt_10k)
```

	Unnamed: 0	region	state	individuals	family_members	\
4	4	Pacific	California	109008.0	20964.0	
9	9	South Atlantic	Florida	21443.0	9587.0	
32	32	Mid-Atlantic	New York	39827.0	52070.0	
37	37	Pacific	Oregon	11139.0	3337.0	
43	43	West South Central	Texas	19199.0	6111.0	
47	47	Pacific	Washington	16424.0	5880.0	

	state_pop
4	39461588
9	21244317
32	19530351
37	4181886
43	28628666
47	7523869

```
In [ ]: # Filter for rows where region is Mountain
mountain_reg = homelessness[homelessness["region"] == "Mountain"]

# See the result
print(mountain_reg)
```

	Unnamed: 0	region	state	individuals	family_members	state_pop
2	2	Mountain	Arizona	7259.0	2606.0	7158024
5	5	Mountain	Colorado	7607.0	3250.0	5691287
12	12	Mountain	Idaho	1297.0	715.0	1750536
26	26	Mountain	Montana	983.0	422.0	1060665
28	28	Mountain	Nevada	7058.0	486.0	3027341
31	31	Mountain	New Mexico	1949.0	602.0	2092741
44	44	Mountain	Utah	1904.0	972.0	3153550
50	50	Mountain	Wyoming	434.0	205.0	577601

```
In [ ]: # Filter for rows where family_members is less than 1000
# and region is Pacific
fam_lt_1k_pac = homelessness[(homelessness["family_members"] < 1000) & (homelessness["region"]=="Pacific")]

# See the result
print(fam_lt_1k_pac)
```

	Unnamed: 0	region	state	individuals	family_members	state_pop
1	1	Pacific	Alaska	1434.0	582.0	735139

Subsetting rows by categorical variables

```
In [ ]: # Subset for rows in South Atlantic or Mid-Atlantic regions
south_mid_atlantic = homelessness[homelessness["region"].isin(["South Atlantic", "Mid-Atlantic"])]

# See the result
print(south_mid_atlantic)
```

	Unnamed: 0	region	state	individuals	\
7	7	South Atlantic	Delaware	708.0	
8	8	South Atlantic	District of Columbia	3770.0	
9	9	South Atlantic	Florida	21443.0	
10	10	South Atlantic	Georgia	6943.0	
20	20	South Atlantic	Maryland	4914.0	
30	30	Mid-Atlantic	New Jersey	6048.0	
32	32	Mid-Atlantic	New York	39827.0	
33	33	South Atlantic	North Carolina	6451.0	
38	38	Mid-Atlantic	Pennsylvania	8163.0	
40	40	South Atlantic	South Carolina	3082.0	
46	46	South Atlantic	Virginia	3928.0	
48	48	South Atlantic	West Virginia	1021.0	

	family_members	state_pop
7	374.0	965479
8	3134.0	701547
9	9587.0	21244317
10	2556.0	10511131
20	2230.0	6035802
30	3350.0	8886025
32	52070.0	19530351
33	2817.0	10381615
38	5349.0	12800922
40	851.0	5084156
46	2047.0	8501286
48	222.0	1804291

```
In [ ]: # The Mojave Desert states
        canu = ["California", "Arizona", "Nevada", "Utah"]

        # Filter for rows in the Mojave Desert states
        mojave_homelessness = homelessness[homelessness["state"].isin(canu)]

        # See the result
        print(mojave_homelessness.head())
```

	Unnamed: 0	region	state	individuals	family_members	state_pop
2	2	Mountain	Arizona	7259.0	2606.0	7158024
4	4	Pacific	California	109008.0	20964.0	39461588
28	28	Mountain	Nevada	7058.0	486.0	3027341
44	44	Mountain	Utah	1904.0	972.0	3153550

Adding new columns

```
In [ ]: # Add total col as sum of individuals and family_members
        homelessness["total"] = homelessness["individuals"] + homelessness["family_members"]
        # Add p_individuals col as proportion of individuals
        homelessness["p_individuals"] = homelessness["individuals"] / homelessness["total"]
```



```
# See the result
print(homelessness.head())
```

Unnamed: 0		region	state	individuals	family_members \
0	0	East South Central	Alabama	2570.0	864.0
1	1	Pacific	Alaska	1434.0	582.0
2	2	Mountain	Arizona	7259.0	2606.0
3	3	West South Central	Arkansas	2280.0	432.0
4	4	Pacific	California	109008.0	20964.0

	state_pop	total	p_individuals
0	4887681	3434.0	0.748398
1	735139	2016.0	0.711310
2	7158024	9865.0	0.735834
3	3009733	2712.0	0.840708
4	39461588	129972.0	0.838704

Combo-attack!

```
In [ ]: # Create indiv_per_10k col as homeless individuals per 10k state pop
homelessness["indiv_per_10k"] = 10000 * homelessness["individuals"] / homelessness["state_pop"]

# Subset rows for indiv_per_10k greater than 20
high_homelessness = homelessness[homelessness["indiv_per_10k"] > 20]

# Sort high_homelessness by descending indiv_per_10k
high_homelessness_srt = high_homelessness.sort_values("indiv_per_10k", ascending=False)

# From high_homelessness_srt, select the state and indiv_per_10k cols
result = high_homelessness_srt[["state", "indiv_per_10k"]]

# See the result
print(result)
```

	state	indiv_per_10k
8	District of Columbia	53.738381
11	Hawaii	29.079406
4	California	27.623825
37	Oregon	26.636307
28	Nevada	23.314189
47	Washington	21.829195
32	New York	20.392363

otros codigos

```
In [ ]: #example using group by
homelessness.groupby("family_members")["individuals"].mean()
print(homelessness.head())
```

	Unnamed: 0	region	state	individuals	family_members	\
0	0	East South Central	Alabama	2570.0	864.0	
1	1	Pacific	Alaska	1434.0	582.0	
2	2	Mountain	Arizona	7259.0	2606.0	
3	3	West South Central	Arkansas	2280.0	432.0	
4	4	Pacific	California	109008.0	20964.0	

	state_pop	total	p_individuals	indiv_per_10k
0	4887681	3434.0	0.748398	5.258117
1	735139	2016.0	0.711310	19.506515
2	7158024	9865.0	0.735834	10.141067
3	3009733	2712.0	0.840708	7.575423
4	39461588	129972.0	0.838704	27.623825

Ch2 Aggregating DataFrames

Aggregating Data

In this chapter, you'll calculate summary statistics on DataFrame columns, and master grouped summary statistics and pivot tables.

```
In [ ]: # Import pandas using the alias pd
import pandas as pd

sales = pd.read_csv('C:\\Users\\yeiso\\OneDrive - Douglas College\\0. DOUGLAS COLLEGE\\3. Fund Machine Learning\\0. Python Course
```

Mean and median

```
In [ ]: # Print the head of the sales DataFrame
print(sales.head())

# Print the info about the sales DataFrame
print(sales.info())

# Print the mean of weekly_sales
print(sales["weekly_sales"].mean())

# Print the median of weekly_sales
print(sales["weekly_sales"].median())
```

	Unnamed: 0	store	type	department	date	weekly_sales	is_holiday	\
0	0	1	A	1	2010-02-05	24924.50	False	
1	1	1	A	1	2010-03-05	21827.90	False	
2	2	1	A	1	2010-04-02	57258.43	False	
3	3	1	A	1	2010-05-07	17413.94	False	
4	4	1	A	1	2010-06-04	17558.09	False	

	temperature_c	fuel_price_usd_per_l	unemployment
0	5.727778	0.679451	8.106
1	8.055556	0.693452	8.106
2	16.816667	0.718284	7.808
3	22.527778	0.748928	7.808
4	27.050000	0.714586	7.808

<class 'pandas.core.frame.DataFrame'>

RangeIndex: 10774 entries, 0 to 10773

Data columns (total 10 columns):

#	Column	Non-Null Count	Dtype
0	Unnamed: 0	10774 non-null	int64
1	store	10774 non-null	int64
2	type	10774 non-null	object
3	department	10774 non-null	int64
4	date	10774 non-null	object
5	weekly_sales	10774 non-null	float64
6	is_holiday	10774 non-null	bool
7	temperature_c	10774 non-null	float64
8	fuel_price_usd_per_l	10774 non-null	float64
9	unemployment	10774 non-null	float64

dtypes: bool(1), float64(4), int64(3), object(2)

memory usage: 768.2+ KB

None

23843.95014850566

12049.064999999999

Summarizing dates

```
In [ ]: # Print the maximum of the date column
print(sales["date"].max())

# Print the minimum of the date column
print(sales["date"].min())
```

2012-10-26

2010-02-05

Efficient summaries

```
In [ ]: # A custom IQR function
def iqr(column):
```

```
return column.quantile(0.75) - column.quantile(0.25)
```

```
# Print IQR of the temperature_c column  
print(sales["temperature_c"].agg(iqr))
```

```
16.583333333333336
```

Cumulative statistics

```
In [ ]: import pandas as pd
```

```
# Sample data
```

```
data = {  
    'store': [5, 1, 4, 9, 8, 7, 10, 3, 1, 6, 11, 2],  
    'type': ['A', 'A', 'A', 'A', 'A', 'A', 'A', 'A', 'A', 'A', 'A', 'A'],  
    'department': [1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1],  
    'date': ['2010-07-02', '2010-02-05', '2010-06-04', '2010-11-05', '2010-10-01', '2010-09-03', '2010-12-03', '2010-05-07', '2010-08-01', '2010-04-01', '2010-03-01', '2010-02-01'],  
    'weekly_sales': [16333.14, 24924.50, 17558.09, 34238.88, 20094.19, 16241.78, 22517.56, 17413.94, 21827.90, 17508.41, 15984.24, 14533.14],  
    'is_holiday': [False, False, False, False, False, False, False, False, False, False, False, False],  
    'temperature_c': [27.172, 5.728, 27.050, 14.856, 22.161, 27.339, 9.594, 22.528, 8.056, 30.644, 9.039, 16.817],  
    'fuel_price_usd_per_l': [0.705, 0.679, 0.715, 0.710, 0.688, 0.681, 0.715, 0.749, 0.693, 0.694, 0.786, 0.718],  
    'unemployment': [7.787, 8.106, 7.808, 7.838, 7.838, 7.787, 7.838, 7.808, 8.106, 7.787, 7.742, 7.808]  
}
```

```
# Create a Pandas DataFrame
```

```
sales_1_1 = pd.DataFrame(data)
```

```
# Display the DataFrame as a table
```

```
sales_1_1
```

Out[]:

	store	type	department	date	weekly_sales	is_holiday	temperature_c	fuel_price_usd_per_l	unemployment
0	5	A	1	2010-07-02	16333.14	False	27.172	0.705	7.787
1	1	A	1	2010-02-05	24924.50	False	5.728	0.679	8.106
2	4	A	1	2010-06-04	17558.09	False	27.050	0.715	7.808
3	9	A	1	2010-11-05	34238.88	False	14.856	0.710	7.838
4	8	A	1	2010-10-01	20094.19	False	22.161	0.688	7.838
5	7	A	1	2010-09-03	16241.78	False	27.339	0.681	7.787
6	10	A	1	2010-12-03	22517.56	False	9.594	0.715	7.838
7	3	A	1	2010-05-07	17413.94	False	22.528	0.749	7.808
8	1	A	1	2010-03-05	21827.90	False	8.056	0.693	8.106
9	6	A	1	2010-08-06	17508.41	False	30.644	0.694	7.787
10	11	A	1	2011-01-07	15984.24	False	9.039	0.786	7.742
11	2	A	1	2010-04-02	57258.43	False	16.817	0.718	7.808

```

In [ ]: # Sort sales_1_1 by date
sales_1_1 = sales_1_1.sort_values(by='date')

# Get the cumulative sum of weekly_sales, add as cum_weekly_sales col
sales_1_1['cum_weekly_sales'] = sales_1_1['weekly_sales'].cumsum()

# Get the cumulative max of weekly_sales, add as cum_max_sales col
sales_1_1['cum_max_sales'] = sales_1_1['weekly_sales'].cummax()

# See the columns you calculated
print(sales_1_1[["date", "weekly_sales", "cum_weekly_sales", "cum_max_sales"]])

```

	date	weekly_sales	cum_weekly_sales	cum_max_sales
1	2010-02-05	24924.50	24924.50	24924.50
8	2010-03-05	21827.90	46752.40	24924.50
11	2010-04-02	57258.43	104010.83	57258.43
7	2010-05-07	17413.94	121424.77	57258.43
2	2010-06-04	17558.09	138982.86	57258.43
0	2010-07-02	16333.14	155316.00	57258.43
9	2010-08-06	17508.41	172824.41	57258.43
5	2010-09-03	16241.78	189066.19	57258.43
4	2010-10-01	20094.19	209160.38	57258.43
3	2010-11-05	34238.88	243399.26	57258.43
6	2010-12-03	22517.56	265916.82	57258.43
10	2011-01-07	15984.24	281901.06	57258.43

Dropping duplicates

```
In [ ]: # Drop duplicate store/type combinations
store_types = sales.drop_duplicates(subset=["store", "type"])
print(store_types.head())

# Drop duplicate store/department combinations
store_depts = sales.drop_duplicates(subset=["store", "department"])
print(store_depts.head())

# Subset the rows where is_holiday is True and drop duplicate dates
holiday_dates = sales[sales["is_holiday"]].drop_duplicates("date")

# Print date col of holiday_dates
print(holiday_dates.head())
```

	Unnamed: 0	store	type	department	date	weekly_sales	\
0	0	1	A	1	2010-02-05	24924.50	
901	901	2	A	1	2010-02-05	35034.06	
1798	1798	4	A	1	2010-02-05	38724.42	
2699	2699	6	A	1	2010-02-05	25619.00	
3593	3593	10	B	1	2010-02-05	40212.84	

	is_holiday	temperature_c	fuel_price_usd_per_l	unemployment
0	False	5.727778	0.679451	8.106
901	False	4.550000	0.679451	8.324
1798	False	6.533333	0.686319	8.623
2699	False	4.683333	0.679451	7.259
3593	False	12.411111	0.782478	9.765

	Unnamed: 0	store	type	department	date	weekly_sales	is_holiday	\
0	0	1	A	1	2010-02-05	24924.50	False	
12	12	1	A	2	2010-02-05	50605.27	False	
24	24	1	A	3	2010-02-05	13740.12	False	
36	36	1	A	4	2010-02-05	39954.04	False	
48	48	1	A	5	2010-02-05	32229.38	False	

	temperature_c	fuel_price_usd_per_l	unemployment
0	5.727778	0.679451	8.106
12	5.727778	0.679451	8.106
24	5.727778	0.679451	8.106
36	5.727778	0.679451	8.106
48	5.727778	0.679451	8.106

	Unnamed: 0	store	type	department	date	weekly_sales	\
498	498	1	A	45	2010-09-10	11.47	
691	691	1	A	77	2011-11-25	1431.00	
2315	2315	4	A	47	2010-02-12	498.00	
6735	6735	19	A	39	2012-09-07	13.41	
6810	6810	19	A	47	2010-12-31	-449.00	

	is_holiday	temperature_c	fuel_price_usd_per_l	unemployment
498	True	25.938889	0.677602	7.787
691	True	15.633333	0.854861	7.866
2315	True	-1.755556	0.679715	8.623
6735	True	22.333333	1.076766	8.193
6810	True	-1.861111	0.881278	8.067

Counting categorical variables

```
In [ ]: # Count the number of stores of each type
store_counts = store_types["type"].value_counts()
print(store_counts)

# Get the proportion of stores of each type
store_props = store_types["type"].value_counts(normalize=True)
print(store_props)
```

```
# Count the number of each department number and sort
dept_counts_sorted = store_depts["department"].value_counts(sort="department", ascending=False)
print(dept_counts_sorted)

# Get the proportion of departments of each number and sort
dept_props_sorted = store_depts["department"].value_counts(sort="department", normalize=True)
print(dept_props_sorted)
```

```
type
A    11
B     1
Name: count, dtype: int64
type
A    0.916667
B    0.083333
Name: proportion, dtype: float64
department
1      12
55     12
72     12
71     12
67     12
...
37     10
48      8
50      6
39      4
43      2
Name: count, Length: 80, dtype: int64
department
1      0.012917
55     0.012917
72     0.012917
71     0.012917
67     0.012917
...
37     0.010764
48     0.008611
50     0.006459
39     0.004306
43     0.002153
Name: proportion, Length: 80, dtype: float64
```

What percent of sales occurred at each store type?

```
In [ ]: # Calc total weekly sales
sales_all = sales["weekly_sales"].sum()

# Subset for type A stores, calc total weekly sales
```



```

sales_A = sales[sales["type"] == "A"]["weekly_sales"].sum()

# Subset for type B stores, calc total weekly sales
sales_B = sales[sales["type"] == "B"]["weekly_sales"].sum()

# Subset for type C stores, calc total weekly sales
sales_C = sales[sales["type"] == "C"]["weekly_sales"].sum()

# Get proportion for each type
sales_propn_by_type = [sales_A, sales_B, sales_C] / sales_all
print(sales_propn_by_type)

```

```
[0.9097747 0.0902253 0.      ]
```

Calculations with .groupby()

```

In [ ]: # Group by type; calc total weekly sales
sales_by_type = sales.groupby("type")["weekly_sales"].sum()

# Get proportion for each type
sales_propn_by_type = sales_by_type / sum(sales_by_type)
print(sales_propn_by_type)

# Group by type and is_holiday; calc total weekly sales
sales_by_type_is_holiday = sales.groupby(["type", "is_holiday"])["weekly_sales"].sum()
print(sales_by_type_is_holiday)

```

```

type
A    0.909775
B    0.090225
Name: weekly_sales, dtype: float64
type  is_holiday
A     False      2.336927e+08
      True       2.360181e+04
B     False      2.317678e+07
      True       1.621410e+03
Name: weekly_sales, dtype: float64

```

Multiple grouped summaries

```

In [ ]: # Import numpy with the alias np
import numpy as np

# For each store type, aggregate weekly_sales: get min, max, mean, and median
sales_stats = sales.groupby("type")["weekly_sales"].agg([min, max, np.mean, np.median])

# Print sales_stats
print(sales_stats)

```

```
# For each store type, aggregate unemployment and fuel_price_usd_per_l: get min, max, mean, and median
unemp_fuel_stats = sales.groupby("type")[["unemployment", "fuel_price_usd_per_l"]].agg([min, max, np.mean, np.median])

# Print unemp_fuel_stats
print(unemp_fuel_stats)
```

	min	max	mean	median	
type					
A	-1098.0	293966.05	23674.667242	11943.92	
B	-798.0	232558.51	25696.678370	13336.08	

	unemployment				fuel_price_usd_per_l		
	min	max	mean	median	min	max	
type							
A	3.879	8.992	7.972611	8.067	0.664129	1.107410	
B	7.170	9.765	9.279323	9.199	0.760023	1.107674	

	mean	median
type		
A	0.744619	0.735455
B	0.805858	0.803348

```
C:\Users\yeiso\AppData\Local\Temp\ipykernel_33652\1976951964.py:5: FutureWarning: The provided callable <built-in function min> is
currently using SeriesGroupBy.min. In a future version of pandas, the provided callable will be used directly. To keep current
behavior pass the string "min" instead.
    sales_stats = sales.groupby("type")["weekly_sales"].agg([min, max, np.mean, np.median])
C:\Users\yeiso\AppData\Local\Temp\ipykernel_33652\1976951964.py:5: FutureWarning: The provided callable <built-in function max> is
currently using SeriesGroupBy.max. In a future version of pandas, the provided callable will be used directly. To keep current
behavior pass the string "max" instead.
    sales_stats = sales.groupby("type")["weekly_sales"].agg([min, max, np.mean, np.median])
C:\Users\yeiso\AppData\Local\Temp\ipykernel_33652\1976951964.py:5: FutureWarning: The provided callable <function mean at 0x00000
1EC3C73B100> is currently using SeriesGroupBy.mean. In a future version of pandas, the provided callable will be used directly. T
o keep current behavior pass the string "mean" instead.
    sales_stats = sales.groupby("type")["weekly_sales"].agg([min, max, np.mean, np.median])
C:\Users\yeiso\AppData\Local\Temp\ipykernel_33652\1976951964.py:5: FutureWarning: The provided callable <function median at 0x000
001EC3C862340> is currently using SeriesGroupBy.median. In a future version of pandas, the provided callable will be used directl
y. To keep current behavior pass the string "median" instead.
    sales_stats = sales.groupby("type")["weekly_sales"].agg([min, max, np.mean, np.median])
C:\Users\yeiso\AppData\Local\Temp\ipykernel_33652\1976951964.py:11: FutureWarning: The provided callable <built-in function min>
is currently using SeriesGroupBy.min. In a future version of pandas, the provided callable will be used directly. To keep current
behavior pass the string "min" instead.
    unemp_fuel_stats = sales.groupby("type")[["unemployment", "fuel_price_usd_per_l"]].agg([min, max, np.mean, np.median])
C:\Users\yeiso\AppData\Local\Temp\ipykernel_33652\1976951964.py:11: FutureWarning: The provided callable <built-in function max>
is currently using SeriesGroupBy.max. In a future version of pandas, the provided callable will be used directly. To keep current
behavior pass the string "max" instead.
    unemp_fuel_stats = sales.groupby("type")[["unemployment", "fuel_price_usd_per_l"]].agg([min, max, np.mean, np.median])
C:\Users\yeiso\AppData\Local\Temp\ipykernel_33652\1976951964.py:11: FutureWarning: The provided callable <function mean at 0x00000
01EC3C73B100> is currently using SeriesGroupBy.mean. In a future version of pandas, the provided callable will be used directly.
To keep current behavior pass the string "mean" instead.
    unemp_fuel_stats = sales.groupby("type")[["unemployment", "fuel_price_usd_per_l"]].agg([min, max, np.mean, np.median])
C:\Users\yeiso\AppData\Local\Temp\ipykernel_33652\1976951964.py:11: FutureWarning: The provided callable <function median at 0x00
0001EC3C862340> is currently using SeriesGroupBy.median. In a future version of pandas, the provided callable will be used direct
ly. To keep current behavior pass the string "median" instead.
    unemp_fuel_stats = sales.groupby("type")[["unemployment", "fuel_price_usd_per_l"]].agg([min, max, np.mean, np.median])
C:\Users\yeiso\AppData\Local\Temp\ipykernel_33652\1976951964.py:11: FutureWarning: The provided callable <built-in function min>
is currently using SeriesGroupBy.min. In a future version of pandas, the provided callable will be used directly. To keep current
behavior pass the string "min" instead.
    unemp_fuel_stats = sales.groupby("type")[["unemployment", "fuel_price_usd_per_l"]].agg([min, max, np.mean, np.median])
```

Pivoting on one variable

```
In [ ]: # Pivot for mean weekly_sales for each store type
mean_sales_by_type = sales.pivot_table(values="weekly_sales", index="type")

# Print mean_sales_by_type
print(mean_sales_by_type)
```

```
      weekly_sales
type
A      23674.667242
B      25696.678370
```

Fill in missing values and sum values with pivot tables

```
In [ ]: # Print mean weekly_sales by department and type; fill missing values with 0
import numpy as np
print(sales.pivot_table(values="weekly_sales", index="department", columns="type", aggfunc=np.mean, fill_value=0))
```

type	A	B
department		
1	30961.725379	44050.626667
2	67600.158788	112958.526667
3	17160.002955	30580.655000
4	44285.399091	51219.654167
5	34821.011364	63236.875000
...
95	123933.787121	77082.102500
96	21367.042857	9528.538333
97	28471.266970	5828.873333
98	12875.423182	217.428333
99	379.123659	0.000000

[80 rows x 2 columns]

C:\Users\yeiso\AppData\Local\Temp\ipykernel_33652\321967142.py:3: FutureWarning: The provided callable <function mean at 0x000001EC3C73B100> is currently using DataFrameGroupBy.mean. In a future version of pandas, the provided callable will be used directly. To keep current behavior pass the string "mean" instead.

```
print(sales.pivot_table(values="weekly_sales", index="department", columns="type", aggfunc=np.mean, fill_value=0))
```

Chapter 3 - Slicing and Indexing DataFrames

Indexes are supercharged row and column names. Learn how they can be combined with slicing for powerful DataFrame subsetting.

```
In [ ]: # Import pandas using the alias pd
import pandas as pd

temperatures = pd.read_csv('C:\\Users\\yeiso\\OneDrive - Douglas College\\0. DOUGLAS COLLEGE\\3. Fund Machine Learning\\0. Python
```

Setting and removing indexes

```
In [ ]: # Look at temperatures
print(temperatures.head())

# Index temperatures by city
temperatures_ind = temperatures.set_index("city")

# Look at temperatures_ind
print(temperatures_ind.head())
```

```
# Reset the index, keeping its contents
print(temperatures_ind.reset_index())

# Reset the index, dropping its contents
print(temperatures_ind.reset_index(drop=True))
```

```
   Unnamed: 0      date      city      country  avg_temp_c
0           0  2000-01-01  Abidjan  Côte D'Ivoire    27.293
1           1  2000-02-01  Abidjan  Côte D'Ivoire    27.685
2           2  2000-03-01  Abidjan  Côte D'Ivoire    29.061
3           3  2000-04-01  Abidjan  Côte D'Ivoire    28.162
4           4  2000-05-01  Abidjan  Côte D'Ivoire    27.547

   Unnamed: 0      date      country  avg_temp_c
city
Abidjan      0  2000-01-01  Côte D'Ivoire    27.293
Abidjan      1  2000-02-01  Côte D'Ivoire    27.685
Abidjan      2  2000-03-01  Côte D'Ivoire    29.061
Abidjan      3  2000-04-01  Côte D'Ivoire    28.162
Abidjan      4  2000-05-01  Côte D'Ivoire    27.547

   city  Unnamed: 0      date      country  avg_temp_c
0  Abidjan      0  2000-01-01  Côte D'Ivoire    27.293
1  Abidjan      1  2000-02-01  Côte D'Ivoire    27.685
2  Abidjan      2  2000-03-01  Côte D'Ivoire    29.061
3  Abidjan      3  2000-04-01  Côte D'Ivoire    28.162
4  Abidjan      4  2000-05-01  Côte D'Ivoire    27.547
...      ...      ...      ...      ...
16495  Xian    16495  2013-05-01      China    18.979
16496  Xian    16496  2013-06-01      China    23.522
16497  Xian    16497  2013-07-01      China    25.251
16498  Xian    16498  2013-08-01      China    24.528
16499  Xian    16499  2013-09-01      China     NaN
```

[16500 rows x 5 columns]

```
   Unnamed: 0      date      country  avg_temp_c
0           0  2000-01-01  Côte D'Ivoire    27.293
1           1  2000-02-01  Côte D'Ivoire    27.685
2           2  2000-03-01  Côte D'Ivoire    29.061
3           3  2000-04-01  Côte D'Ivoire    28.162
4           4  2000-05-01  Côte D'Ivoire    27.547
...      ...      ...      ...
16495    16495  2013-05-01      China    18.979
16496    16496  2013-06-01      China    23.522
16497    16497  2013-07-01      China    25.251
16498    16498  2013-08-01      China    24.528
16499    16499  2013-09-01      China     NaN
```

[16500 rows x 4 columns]

Subsetting with .loc[]

```
In [ ]: # Make a list of cities to subset on
cities = ["Moscow", "Saint Petersburg"]

# Subset temperatures using square brackets
print(temperatures[temperatures["city"].isin(cities)])

# Subset temperatures_ind using .loc[]
print(temperatures_ind.loc[cities])
```

	Unnamed: 0	date	city	country	avg_temp_c
10725	10725	2000-01-01	Moscow	Russia	-7.313
10726	10726	2000-02-01	Moscow	Russia	-3.551
10727	10727	2000-03-01	Moscow	Russia	-1.661
10728	10728	2000-04-01	Moscow	Russia	10.096
10729	10729	2000-05-01	Moscow	Russia	10.357
...
13360	13360	2013-05-01	Saint Petersburg	Russia	12.355
13361	13361	2013-06-01	Saint Petersburg	Russia	17.185
13362	13362	2013-07-01	Saint Petersburg	Russia	17.234
13363	13363	2013-08-01	Saint Petersburg	Russia	17.153
13364	13364	2013-09-01	Saint Petersburg	Russia	NaN

[330 rows x 5 columns]

	Unnamed: 0	date	country	avg_temp_c
city				
Moscow	10725	2000-01-01	Russia	-7.313
Moscow	10726	2000-02-01	Russia	-3.551
Moscow	10727	2000-03-01	Russia	-1.661
Moscow	10728	2000-04-01	Russia	10.096
Moscow	10729	2000-05-01	Russia	10.357
...
Saint Petersburg	13360	2013-05-01	Russia	12.355
Saint Petersburg	13361	2013-06-01	Russia	17.185
Saint Petersburg	13362	2013-07-01	Russia	17.234
Saint Petersburg	13363	2013-08-01	Russia	17.153
Saint Petersburg	13364	2013-09-01	Russia	NaN

[330 rows x 4 columns]

Setting multi-level indexes

```
In [ ]: # Index temperatures by country & city
temperatures_ind = temperatures.set_index(["country", "city"])

# List of tuples: Brazil, Rio De Janeiro & Pakistan, Lahore
rows_to_keep = [("Brazil", "Rio De Janeiro"), ("Pakistan", "Lahore")]
```

```
# Subset for rows to keep
print(temperatures_ind.loc[rows_to_keep])
```

		Unnamed: 0	date	avg_temp_c
country	city			
Brazil	Rio De Janeiro	12540	2000-01-01	25.974
	Rio De Janeiro	12541	2000-02-01	26.699
	Rio De Janeiro	12542	2000-03-01	26.270
	Rio De Janeiro	12543	2000-04-01	25.750
	Rio De Janeiro	12544	2000-05-01	24.356
...	
Pakistan	Lahore	8575	2013-05-01	33.457
	Lahore	8576	2013-06-01	34.456
	Lahore	8577	2013-07-01	33.279
	Lahore	8578	2013-08-01	31.511
	Lahore	8579	2013-09-01	NaN

[330 rows x 5 columns]

Sorting by index values

```
In [ ]: # Sort temperatures_ind by index values
print(temperatures_ind.sort_index())

# Sort temperatures_ind by index values at the city level
print(temperatures_ind.sort_index(level=["city", "country"]))

# Sort temperatures_ind by country then descending city
print(temperatures_ind.sort_index(level=["country", "city"], ascending=[True, False]))
```

		Unnamed: 0	date	avg_temp_c
Afghanistan	Kabul	7260	2000-01-01	3.326
	Kabul	7261	2000-02-01	3.454
	Kabul	7262	2000-03-01	9.612
	Kabul	7263	2000-04-01	17.925
	Kabul	7264	2000-05-01	24.658
...	
Zimbabwe	Harare	5605	2013-05-01	18.298
	Harare	5606	2013-06-01	17.020
	Harare	5607	2013-07-01	16.299
	Harare	5608	2013-08-01	19.232
	Harare	5609	2013-09-01	NaN

[16500 rows x 3 columns]

		Unnamed: 0	date	avg_temp_c
Côte D'Ivoire	Abidjan	0	2000-01-01	27.293
	Abidjan	1	2000-02-01	27.685
	Abidjan	2	2000-03-01	29.061
	Abidjan	3	2000-04-01	28.162
	Abidjan	4	2000-05-01	27.547
...	
China	Xian	16495	2013-05-01	18.979
	Xian	16496	2013-06-01	23.522
	Xian	16497	2013-07-01	25.251
	Xian	16498	2013-08-01	24.528
	Xian	16499	2013-09-01	NaN

[16500 rows x 3 columns]

		Unnamed: 0	date	avg_temp_c
Afghanistan	Kabul	7260	2000-01-01	3.326
	Kabul	7261	2000-02-01	3.454
	Kabul	7262	2000-03-01	9.612
	Kabul	7263	2000-04-01	17.925
	Kabul	7264	2000-05-01	24.658
...	
Zimbabwe	Harare	5605	2013-05-01	18.298
	Harare	5606	2013-06-01	17.020
	Harare	5607	2013-07-01	16.299
	Harare	5608	2013-08-01	19.232
	Harare	5609	2013-09-01	NaN

[16500 rows x 3 columns]

Slicing index values


```
In [ ]: # Sort the index of temperatures_ind
temperatures_srt = temperatures_ind.sort_index()

# Subset rows from Pakistan to Russia
print(temperatures_srt.loc["Pakistan":"Russia"])

# Try to subset rows from Lahore to Moscow
print(temperatures_srt.loc["Lahore":"Moscow"])

# Subset rows from Pakistan, Lahore to Russia, Moscow
print(temperatures_srt.loc[("Pakistan","Lahore"):(("Russia","Moscow")])
```

		Unnamed: 0	date	avg_temp_c
country	city			
Pakistan	Faisalabad	4785	2000-01-01	12.792
	Faisalabad	4786	2000-02-01	14.339
	Faisalabad	4787	2000-03-01	20.309
	Faisalabad	4788	2000-04-01	29.072
	Faisalabad	4789	2000-05-01	34.845
...	
Russia	Saint Petersburg	13360	2013-05-01	12.355
	Saint Petersburg	13361	2013-06-01	17.185
	Saint Petersburg	13362	2013-07-01	17.234
	Saint Petersburg	13363	2013-08-01	17.153
	Saint Petersburg	13364	2013-09-01	NaN

[1155 rows x 3 columns]

		Unnamed: 0	date	avg_temp_c
country	city			
Mexico	Mexico	10230	2000-01-01	12.694
	Mexico	10231	2000-02-01	14.677
	Mexico	10232	2000-03-01	17.376
	Mexico	10233	2000-04-01	18.294
	Mexico	10234	2000-05-01	18.562
...	
Morocco	Casablanca	3130	2013-05-01	19.217
	Casablanca	3131	2013-06-01	23.649
	Casablanca	3132	2013-07-01	27.488
	Casablanca	3133	2013-08-01	27.952
	Casablanca	3134	2013-09-01	NaN

[330 rows x 3 columns]

		Unnamed: 0	date	avg_temp_c
country	city			
Pakistan	Lahore	8415	2000-01-01	12.792
	Lahore	8416	2000-02-01	14.339
	Lahore	8417	2000-03-01	20.309
	Lahore	8418	2000-04-01	29.072
	Lahore	8419	2000-05-01	34.845
...	
Russia	Moscow	10885	2013-05-01	16.152
	Moscow	10886	2013-06-01	18.718
	Moscow	10887	2013-07-01	18.136
	Moscow	10888	2013-08-01	17.485
	Moscow	10889	2013-09-01	NaN

[660 rows x 3 columns]

Slicing in both directions

```
In [ ]: # Subset rows from India, Hyderabad to Iraq, Baghdad
print(temperatures_srt.loc[("India", "Hyderabad"):(("Iraq", "Baghdad"))])

# Subset columns from date to avg_temp_c
print(temperatures_srt.loc[:, "date": "avg_temp_c"])

# Subset in both directions at once
print(temperatures_srt.loc[("India", "Hyderabad"):(("Iraq", "Baghdad"), "date": "avg_temp_c"])
```

		Unnamed: 0	date	avg_temp_c
country	city			
India	Hyderabad	5940	2000-01-01	23.779
	Hyderabad	5941	2000-02-01	25.826
	Hyderabad	5942	2000-03-01	28.821
	Hyderabad	5943	2000-04-01	32.698
	Hyderabad	5944	2000-05-01	32.438
...	
Iraq	Baghdad	1150	2013-05-01	28.673
	Baghdad	1151	2013-06-01	33.803
	Baghdad	1152	2013-07-01	36.392
	Baghdad	1153	2013-08-01	35.463
	Baghdad	1154	2013-09-01	NaN

[2145 rows x 3 columns]

		date	avg_temp_c
country	city		
Afghanistan	Kabul	2000-01-01	3.326
	Kabul	2000-02-01	3.454
	Kabul	2000-03-01	9.612
	Kabul	2000-04-01	17.925
	Kabul	2000-05-01	24.658
...	
Zimbabwe	Harare	2013-05-01	18.298
	Harare	2013-06-01	17.020
	Harare	2013-07-01	16.299
	Harare	2013-08-01	19.232
	Harare	2013-09-01	NaN

[16500 rows x 2 columns]

		date	avg_temp_c
country	city		
India	Hyderabad	2000-01-01	23.779
	Hyderabad	2000-02-01	25.826
	Hyderabad	2000-03-01	28.821
	Hyderabad	2000-04-01	32.698
	Hyderabad	2000-05-01	32.438
...	
Iraq	Baghdad	2013-05-01	28.673
	Baghdad	2013-06-01	33.803
	Baghdad	2013-07-01	36.392
	Baghdad	2013-08-01	35.463
	Baghdad	2013-09-01	NaN

[2145 rows x 2 columns]

Slicing time series

```
In [ ]: # Use Boolean conditions to subset temperatures for rows in 2010 and 2011
temperatures_bool = temperatures[(temperatures["date"] >= "2010-01-01") & (temperatures["date"] <= "2011-12-31")]
print(temperatures_bool)

# Set date as an index and sort the index
temperatures_ind = temperatures.set_index("date").sort_index()

# Use .loc[] to subset temperatures_ind for rows in 2010 and 2011
print(temperatures_ind.loc["2010":"2011"])

# Use .loc[] to subset temperatures_ind for rows from Aug 2010 to Feb 2011
print(temperatures_ind.loc["2010-08":"2011-02"])
```

	Unnamed: 0	date	city	country	avg_temp_c
120	120	2010-01-01	Abidjan	Côte D'Ivoire	28.270
121	121	2010-02-01	Abidjan	Côte D'Ivoire	29.262
122	122	2010-03-01	Abidjan	Côte D'Ivoire	29.596
123	123	2010-04-01	Abidjan	Côte D'Ivoire	29.068
124	124	2010-05-01	Abidjan	Côte D'Ivoire	28.258
...
16474	16474	2011-08-01	Xian	China	23.069
16475	16475	2011-09-01	Xian	China	16.775
16476	16476	2011-10-01	Xian	China	12.587
16477	16477	2011-11-01	Xian	China	7.543
16478	16478	2011-12-01	Xian	China	-0.490

[2400 rows x 5 columns]

	Unnamed: 0	city	country	avg_temp_c
date				
2010-01-01	4905	Faisalabad	Pakistan	11.810
2010-01-01	10185	Melbourne	Australia	20.016
2010-01-01	3750	Chongqing	China	7.921
2010-01-01	13155	São Paulo	Brazil	23.738
2010-01-01	5400	Guangzhou	China	14.136
...
2010-12-01	6896	Jakarta	Indonesia	26.602
2010-12-01	5246	Gizeh	Egypt	16.530
2010-12-01	11186	Nagpur	India	19.120
2010-12-01	14981	Sydney	Australia	19.559
2010-12-01	13496	Salvador	Brazil	26.265

[1200 rows x 4 columns]

	Unnamed: 0	city	country	avg_temp_c
date				
2010-08-01	2602	Calcutta	India	30.226
2010-08-01	12337	Pune	India	24.941
2010-08-01	6562	Izmir	Turkey	28.352
2010-08-01	15637	Tianjin	China	25.543
2010-08-01	9862	Manila	Philippines	27.101
...
2011-01-01	4257	Dar Es Salaam	Tanzania	28.541
2011-01-01	11352	Nairobi	Kenya	17.768
2011-01-01	297	Addis Abeba	Ethiopia	17.708
2011-01-01	11517	Nanjing	China	0.144
2011-01-01	11847	New York	United States	-4.463

[600 rows x 4 columns]

Subsetting by row/column number

```
In [ ]: # Get 23rd row, 2nd column (index 22, 1)
print(temperatures.iloc[22,1])
```

```
# Use slicing to get the first 5 rows
print(temperatures.iloc[:5])

# Use slicing to get columns 3 to 4
print(temperatures.iloc[:,2:4])

# Use slicing in both directions at once
print(temperatures.iloc[:5,2:4])
```

```
2001-11-01
   Unnamed: 0      date      city      country  avg_temp_c
0           0  2000-01-01  Abidjan  Côte D'Ivoire      27.293
1           1  2000-02-01  Abidjan  Côte D'Ivoire      27.685
2           2  2000-03-01  Abidjan  Côte D'Ivoire      29.061
3           3  2000-04-01  Abidjan  Côte D'Ivoire      28.162
4           4  2000-05-01  Abidjan  Côte D'Ivoire      27.547
   city      country
0  Abidjan  Côte D'Ivoire
1  Abidjan  Côte D'Ivoire
2  Abidjan  Côte D'Ivoire
3  Abidjan  Côte D'Ivoire
4  Abidjan  Côte D'Ivoire
...      ...      ...
16495    Xian      China
16496    Xian      China
16497    Xian      China
16498    Xian      China
16499    Xian      China
```

```
[16500 rows x 2 columns]
```

```
   city      country
0  Abidjan  Côte D'Ivoire
1  Abidjan  Côte D'Ivoire
2  Abidjan  Côte D'Ivoire
3  Abidjan  Côte D'Ivoire
4  Abidjan  Côte D'Ivoire
```

Pivot temperature by city and year

```
In [ ]: import pandas as pd

# Assuming you have a DataFrame named temperatures

# Convert "date" column to datetime format
temperatures["date"] = pd.to_datetime(temperatures["date"])

# Add a year column to temperatures
temperatures["year"] = temperatures["date"].dt.year
```

```
# Pivot avg_temp_c by country and city vs year
```

```
temp_by_country_city_vs_year = temperatures.pivot_table(values="avg_temp_c", index=["country", "city"], columns="year")
```

```
# See the result
```

```
print(temp_by_country_city_vs_year)
```


year		2000	2001	2002	2003 \
country	city				
Afghanistan	Kabul	15.822667	15.847917	15.714583	15.132583
Angola	Luanda	24.410333	24.427083	24.790917	24.867167
Australia	Melbourne	14.320083	14.180000	14.075833	13.985583
	Sydney	17.567417	17.854500	17.733833	17.592333
Bangladesh	Dhaka	25.905250	25.931250	26.095000	25.927417
...	
United States	Chicago	11.089667	11.703083	11.532083	10.481583
	Los Angeles	16.643333	16.466250	16.430250	16.944667
	New York	9.969083	10.931000	11.252167	9.836000
Vietnam	Ho Chi Minh City	27.588917	27.831750	28.064750	27.827667
Zimbabwe	Harare	20.283667	20.861000	21.079333	20.889167

year		2004	2005	2006	2007 \
country	city				
Afghanistan	Kabul	16.128417	14.847500	15.798500	15.518000
Angola	Luanda	24.216167	24.414583	24.138417	24.241583
Australia	Melbourne	13.742083	14.378500	13.991083	14.991833
	Sydney	17.869667	18.028083	17.749500	18.020833
Bangladesh	Dhaka	26.136083	26.193333	26.440417	25.951333
...	
United States	Chicago	10.943417	11.583833	11.870500	11.448333
	Los Angeles	16.552833	16.431417	16.623083	16.699917
	New York	10.389500	10.681417	11.519250	10.627333
Vietnam	Ho Chi Minh City	27.686583	27.884000	28.044000	27.866667
Zimbabwe	Harare	20.307667	21.487417	20.699750	20.746250

year		2008	2009	2010	2011 \
country	city				
Afghanistan	Kabul	15.479250	15.093333	15.676000	15.812167
Angola	Luanda	24.266333	24.325083	24.440250	24.150750
Australia	Melbourne	14.110583	14.647417	14.231667	14.190917
	Sydney	17.321083	18.175833	17.999000	17.713333
Bangladesh	Dhaka	26.004500	26.535583	26.648167	25.803250
...	
United States	Chicago	10.242417	10.298333	11.815917	11.214250
	Los Angeles	17.014750	16.677000	15.887000	15.874833
	New York	10.641667	10.141833	11.357583	11.272250
Vietnam	Ho Chi Minh City	27.611417	27.853333	28.281750	27.675417
Zimbabwe	Harare	20.680500	20.523833	21.165833	20.781750

year		2012	2013
country	city		
Afghanistan	Kabul	14.510333	16.206125
Angola	Luanda	24.240083	24.553875
Australia	Melbourne	14.268667	14.741500
	Sydney	17.474333	18.089750
Bangladesh	Dhaka	26.283583	26.587000
...	

United States	Chicago	12.821250	11.586889
	Los Angeles	17.089583	18.120667
	New York	11.971500	12.163889
Vietnam	Ho Chi Minh City	28.248750	28.455000
Zimbabwe	Harare	20.523333	19.756500

[100 rows x 4 columns]

Subsetting pivot tables

```
In [ ]: # Subset for Egypt to India
temp_by_country_city_vs_year.loc["Egypt":"India"]

# Subset for Egypt, Cairo to India, Delhi
temp_by_country_city_vs_year.loc[("Egypt", "Cairo"):(("India", "Delhi"))]

# Subset in both directions at once
temp_by_country_city_vs_year.loc[("Egypt", "Cairo"):(("India", "Delhi"), "2005":"2010")]
```

```
Out [ ]:
```

	year	2005	2006	2007	2008	2009	2010
country	city						
Egypt	Cairo	22.006500	22.050000	22.361000	22.644500	22.625000	23.718250
	Gizeh	22.006500	22.050000	22.361000	22.644500	22.625000	23.718250
Ethiopia	Addis Abeba	18.312833	18.427083	18.142583	18.165000	18.765333	18.298250
France	Paris	11.552917	11.788500	11.750833	11.278250	11.464083	10.409833
Germany	Berlin	9.919083	10.545333	10.883167	10.657750	10.062500	8.606833
India	Ahmadabad	26.828083	27.282833	27.511167	27.048500	28.095833	28.017833
	Bangalore	25.476500	25.418250	25.464333	25.352583	25.725750	25.705250
	Bombay	27.035750	27.381500	27.634667	27.177750	27.844500	27.765417
	Calcutta	26.729167	26.986250	26.584583	26.522333	27.153250	27.288833
	Delhi	25.716083	26.365917	26.145667	25.675000	26.554250	26.520250

Calculating on a pivot table

```
In [ ]: # Get the worldwide mean temp by year
mean_temp_by_year = temp_by_country_city_vs_year.mean(axis="index")

# Filter for the year that had the highest mean temp
print(mean_temp_by_year[mean_temp_by_year == max(mean_temp_by_year)])
```

```
# Get the mean temp by city
mean_temp_by_city = temp_by_country_city_vs_year.mean(axis="columns")

# Filter for the city that had the lowest mean temp
print(mean_temp_by_city[mean_temp_by_city == min(mean_temp_by_city)])

year
2013    20.312285
dtype: float64
country city
China   Harbin    4.876551
dtype: float64
```

hchapter 4 - Creating and Visualizing DataFrames

Learn to visualize the contents of your DataFrames, handle missing data values, and import data from and export data to CSV files.

```
In [ ]: # Import pandas using the alias pd
import pandas as pd

avocados = pd.read_pickle('C:\\Users\\yeiso\\OneDrive - Douglas College\\0. DOUGLAS COLLEGE\\3. Fund Machine Learning\\0. Python C
print(avocados.head(3))
```

	date	type	year	avg_price	size	nb_sold
0	2015-12-27	conventional	2015	0.95	small	9626901.09
1	2015-12-20	conventional	2015	0.98	small	8710021.76
2	2015-12-13	conventional	2015	0.93	small	9855053.66

Which avocado size is most popular?

```
In [ ]: # Import matplotlib.pyplot with alias plt
import matplotlib.pyplot as plt


# Look at the first few rows of data
print(avocados.head())

# Get the total number of avocados sold of each size
nb_sold_by_size = avocados.groupby("size")["nb_sold"].sum()

# Create a bar plot of the number of avocados sold by size
nb_sold_by_size.plot(kind="bar")

# Show the plot
plt.show()
```

	date	type	year	avg_price	size	nb_sold
0	2015-12-27	conventional	2015	0.95	small	9626901.09
1	2015-12-20	conventional	2015	0.98	small	8710021.76
2	2015-12-13	conventional	2015	0.93	small	9855053.66
3	2015-12-06	conventional	2015	0.89	small	9405464.36
4	2015-11-29	conventional	2015	0.99	small	8094803.56




Changes in sales over time

```
In [ ]: # Import matplotlib.pyplot with alias plt
import matplotlib.pyplot as plt

# Get the total number of avocados sold on each date
nb_sold_by_date = avocados.groupby("date")["nb_sold"].sum()

# Create a line plot of the number of avocados sold by date
nb_sold_by_date.plot(kind="line")
```


```
# Show the plot  
plt.show()
```



Avocado supply and demand

```
In [ ]: # Scatter plot of nb_sold vs avg_price with title  
avocados.plot(x="nb_sold", y="avg_price", kind="scatter", title="Number of avocados sold vs. average price")  
  
# Show the plot  
plt.show()
```

Number of avocados sold vs. average price




Price of conventional vs. organic avocados

```
In [ ]: # Histogram of conventional avg_price
avocados[avocados["type"] == "conventional"]["avg_price"].hist()

# Histogram of organic avg_price
avocados[avocados["type"] == "organic"]["avg_price"].hist()

# Add a Legend
plt.legend(["conventional", "organic"])

# Show the plot
plt.show()
```



Finding missing values

```
In [ ]: # Import matplotlib.pyplot with alias plt
import matplotlib.pyplot as plt

# Check individual values for missing values
print(avocados_2016.isna())

# Check each column for missing values
print(avocados_2016.isna().any())

# Bar plot of missing values by variable
avocados_2016.isna().sum().plot(kind="bar")

# Show plot
plt.show()
```

Removing missing values

```
In [ ]: # Remove rows with missing values
avocados_complete = avocados_2016.dropna()
```

```
# Check if any columns contain missing values
print(avocados_complete.isna().any())
```

Replacing missing values

```
In [ ]: # List the columns with missing values
cols_with_missing = ["small_sold", "large_sold", "xl_sold"]

# Create histograms showing the distributions cols_with_missing
avocados_2016[cols_with_missing].plot(kind="hist")

# Fill in missing values with 0
avocados_filled = avocados_2016.fillna(0)

# Create histograms of the filled columns
avocados_filled[cols_with_missing].hist()

# Show the plot
plt.show()
```

List of dictionaries

```
In [ ]: # Create a list of dictionaries with new data
avocados_list = [
    {"date": "2019-11-03", "small_sold": 10376832, "large_sold": 7835071},
    {"date": "2019-11-10", "small_sold": 10717154, "large_sold": 8561348},
]

# Convert list into DataFrame
avocados_2019 = pd.DataFrame(avocados_list)

# Print the new DataFrame
print(avocados_2019)
```

	date	small_sold	large_sold
0	2019-11-03	10376832	7835071
1	2019-11-10	10717154	8561348

Dictionary of lists

```
In [ ]: # Create a dictionary of lists with new data
avocados_dict = {
    "date": ["2019-11-17", "2019-12-01"],
    "small_sold": [10859987, 9291631],
    "large_sold": [7674135, 6238096]
}

# Convert dictionary into DataFrame
```



```
avocados_2019 = pd.DataFrame(avocados_dict)
```

```
# Print the new DataFrame
```

```
print(avocados_2019)
```

```
      date  small_sold  large_sold
0  2019-11-17    10859987    7674135
1  2019-12-01     9291631    6238096
```

CSV to DataFrame

```
In [ ]: # Read CSV as DataFrame called airline_bumping
airline_bumping = pd.read_csv("airline_bumping.csv")

# Take a Look at the DataFrame
print(airline_bumping.head())

# For each airline, select nb_bumped and total_passengers and sum
airline_totals = airline_bumping.groupby("airline")[["nb_bumped", "total_passengers"]].sum()

# Create new col, bumps_per_10k: no. of bumps per 10k passengers for each airline
airline_totals["bumps_per_10k"] = airline_totals["nb_bumped"] / airline_totals["total_passengers"] * 10000

# Print airline_totals
print(airline_totals)
```

DataFrame to CSV

```
In [ ]: # Create airline_totals_sorted
airline_totals_sorted = airline_totals.sort_values("bumps_per_10k", ascending=False)

# Print airline_totals_sorted
print(airline_totals_sorted)

# Save as airline_totals_sorted.csv
airline_totals_sorted.to_csv("airline_totals_sorted.csv")
```