

9. Manipulating Time Series Data in Python

Chapter 1 - Working with Time Series in Pandas

In this course you'll learn the basics of manipulating time series data. Time series data are data that are indexed by a sequence of dates or times. You'll learn how to use methods built into Pandas to work with this index. You'll also learn how resample time series to change the frequency. This course will also show you how to calculate rolling and cumulative values for times series. Finally, you'll use all your new skills to build a value-weighted stock index from actual stock data.

```
In [ ]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns

plt.rcParams['figure.figsize'] = (10, 5)
```

Your first time series

```
In [ ]: seven_days = pd.date_range(start='2017-1-1', periods=7)

# Iterate over the dates and print the number and name of the weekday
for day in seven_days:
    print(day.dayofweek, day.day_name())
```

```
6 Sunday
0 Monday
1 Tuesday
2 Wednesday
3 Thursday
4 Friday
5 Saturday
```

Create a time series of air quality data

```
In [ ]: #reading file "NYC"
data = pd.read_csv('C:\\Users\\yeiso\\OneDrive - Douglas College\\0. DOUGLAS COLLEGE\\3. Fund Machine Learning\\0. Python Course D

# Inspect data
print(data.info())
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 6317 entries, 0 to 6316
Data columns (total 4 columns):
#   Column  Non-Null Count  Dtype
---  -
0    date    6317 non-null    object
1   ozone    6317 non-null    float64
2   pm25     6317 non-null    float64
3    co      6317 non-null    float64
dtypes: float64(3), object(1)
memory usage: 197.5+ KB
None
```


```
In [ ]: # Convert the date column to datetime64
data['date'] = pd.to_datetime(data['date'])
```

```
In [ ]: # Set date column as index
data.set_index('date', inplace=True)
```

```
In [ ]: # Inspect data
print(data.info())
```

```
# Plot data
data.plot(subplots=True);
plt.show()
```

```
<class 'pandas.core.frame.DataFrame'>
DatetimeIndex: 6317 entries, 1999-07-01 to 2017-03-31
Data columns (total 3 columns):
#   Column  Non-Null Count  Dtype
---  -
0   ozone    6317 non-null    float64
1   pm25     6317 non-null    float64
2    co      6317 non-null    float64
dtypes: float64(3)
memory usage: 197.4 KB
None
```



Compare annual stock price trends


```
In [ ]: df = pd.read_csv('C:\\Users\\yeiso\\OneDrive - Douglas College\\0. DOUGLAS COLLEGE\\3. Fund Machine Learning\\0. Python Course Data\\stock_prices.csv')
#df.drop(['Unnamed: 0'], axis=1, inplace=True)
df.head()
```

```
Out[ ]:      price
date
2013-01-02  20.08
2013-01-03  19.78
2013-01-04  19.86
2013-01-07  19.40
2013-01-08  19.66
```

```
In [ ]: yahoo = df
# Create dataframe prices here
prices = pd.DataFrame()
```

```
# Select data for each year and concatenate with prices here
for year in ['2013', '2014', '2015']:
    price_per_year = yahoo.loc[year, ['price']].reset_index(drop=True)
    price_per_year.rename(columns={'price': year}, inplace=True)
    prices = pd.concat([prices, price_per_year], axis=1)

prices.plot()
plt.show()
```



```
In [ ]: print(prices.head())
```

	2013	2014	2015
0	20.08	NaN	NaN
1	19.78	39.59	50.17
2	19.86	40.12	49.13
3	19.40	39.93	49.21
4	19.66	40.92	48.59

Set and change time series frequency

```
In [ ]: co = pd.read_csv('C:\\Users\\yeiso\\OneDrive - Douglas College\\0. DOUGLAS COLLEGE\\3. Fund Machine Learning\\0. Python Course Dat
co['date'] = pd.to_datetime(co['date'])
co.set_index('date', inplace=True)
co.head()
```

```
Out[ ]:           Chicago  Los Angeles  New York
```

date			
2005-01-01	0.317763	0.777657	0.639830
2005-01-03	0.520833	0.349547	0.969572
2005-01-04	0.477083	0.626630	0.905208
2005-01-05	0.348822	0.613814	0.769176
2005-01-06	0.572917	0.792596	0.815761

```
In [ ]: print(co.info())


# Set the frequency to calendar daily
co = co.asfreq('D')


# Plot the data
co.plot(subplots=True);
plt.show()

# Set Frequency to monthly
co = co.asfreq('M')

# Plot the data
co.plot(subplots=True)
plt.show()

<class 'pandas.core.frame.DataFrame'>
DatetimeIndex: 1898 entries, 2005-01-01 to 2010-12-31
Data columns (total 3 columns):
#   Column          Non-Null Count  Dtype
---  -
0    Chicago         1898 non-null   float64
1    Los Angeles     1898 non-null   float64
2    New York        1898 non-null   float64
dtypes: float64(3)
memory usage: 59.3 KB
None
```






Shifting stock prices across time

```
In [ ]: google = pd.read_csv('C:\\Users\\yeiso\\OneDrive - Douglas College\\0. DOUGLAS COLLEGE\\3. Fund Machine Learning\\0. Python Course
# Set data frequency to business daily
google = google.asfreq('B')

# Create 'lagged' and 'shifted'
google['lagged'] = google['Close'].shift(periods=-90)
google['shifted'] = google['Close'].shift(periods=90)

# Plot the google price series
google.plot();
plt.savefig('C:\\Users\\yeiso\\OneDrive - Douglas College\\0. DOUGLAS COLLEGE\\3. Fund Machine Learning\\0. Python Course DataCamp
```

```
c:\Users\yeiso\AppData\Local\Programs\Python\Python312\Lib\site-packages\matplotlib\axis.py:1769: FutureWarning: Period with BDay
freq is deprecated and will be removed in a future version. Use a DatetimeIndex with BDay freq instead.
    ret = self.converter.convert(x, self.units, self)
c:\Users\yeiso\AppData\Local\Programs\Python\Python312\Lib\site-packages\matplotlib\axis.py:1769: FutureWarning: PeriodDtype[B] i
s deprecated and will be removed in a future version. Use a DatetimeIndex with freq='B' instead
    ret = self.converter.convert(x, self.units, self)
c:\Users\yeiso\AppData\Local\Programs\Python\Python312\Lib\site-packages\matplotlib\axis.py:1769: FutureWarning: PeriodDtype[B] i
s deprecated and will be removed in a future version. Use a DatetimeIndex with freq='B' instead
    ret = self.converter.convert(x, self.units, self)
c:\Users\yeiso\AppData\Local\Programs\Python\Python312\Lib\site-packages\matplotlib\axis.py:1769: FutureWarning: PeriodDtype[B] i
s deprecated and will be removed in a future version. Use a DatetimeIndex with freq='B' instead
    ret = self.converter.convert(x, self.units, self)
c:\Users\yeiso\AppData\Local\Programs\Python\Python312\Lib\site-packages\matplotlib\axis.py:1495: FutureWarning: Period with BDay
freq is deprecated and will be removed in a future version. Use a DatetimeIndex with BDay freq instead.
    return self.major.locator()
c:\Users\yeiso\AppData\Local\Programs\Python\Python312\Lib\site-packages\matplotlib\axis.py:1495: FutureWarning: PeriodDtype[B] i
s deprecated and will be removed in a future version. Use a DatetimeIndex with freq='B' instead
    return self.major.locator()
c:\Users\yeiso\AppData\Local\Programs\Python\Python312\Lib\site-packages\matplotlib\ticker.py:216: FutureWarning: Period with BDa
y freq is deprecated and will be removed in a future version. Use a DatetimeIndex with BDay freq instead.
    return [self(value, i) for i, value in enumerate(values)]
```

Calculating stock price changes

```
In [ ]: yahoo = yahoo.asfreq('B')
```

```
In [ ]: yahoo['shifted_30'] = yahoo['price'].shift(periods=30)

# Subtract shifted_30 from price
yahoo['change_30'] = yahoo['price'] - yahoo['shifted_30']

# Get the 30-day price difference
yahoo['diff_30'] = yahoo['price'].diff(periods=30)

# Inspect the last five rows of price
print(yahoo['price'].tail(5))

# Show the value_counts of the difference between change_30 and diff_30
print(yahoo['diff_30'].sub(yahoo['change_30']).value_counts())
```

```
date
2015-12-25      NaN
2015-12-28     33.60
2015-12-29     34.04
2015-12-30     33.37
2015-12-31     33.26
Freq: B, Name: price, dtype: float64
0.0      703
Name: count, dtype: int64
```

Plotting multi-period returns

```
In [ ]: google = pd.read_csv('C:\\Users\\yeiso\\OneDrive - Douglas College\\0. DOUGLAS COLLEGE\\3. Fund Machine Learning\\0. Python Course

# Set data frequency to business daily
google = google.asfreq('D')
```

```
In [ ]: google['daily_return'] = google['Close'].pct_change(periods=1) * 100

# Create monthly_return
google['monthly_return'] = google['Close'].pct_change(periods=30) * 100

# Create annual_return
google['annual_return'] = google['Close'].pct_change(periods=360) * 100

# Plot the result
google.plot(subplots=True);
```

C:\Users\yeiso\AppData\Local\Temp\ipykernel_19956\207313826.py:1: FutureWarning: The default fill_method='pad' in Series.pct_change is deprecated and will be removed in a future version. Call ffill before calling pct_change to retain current behavior and silence this warning.

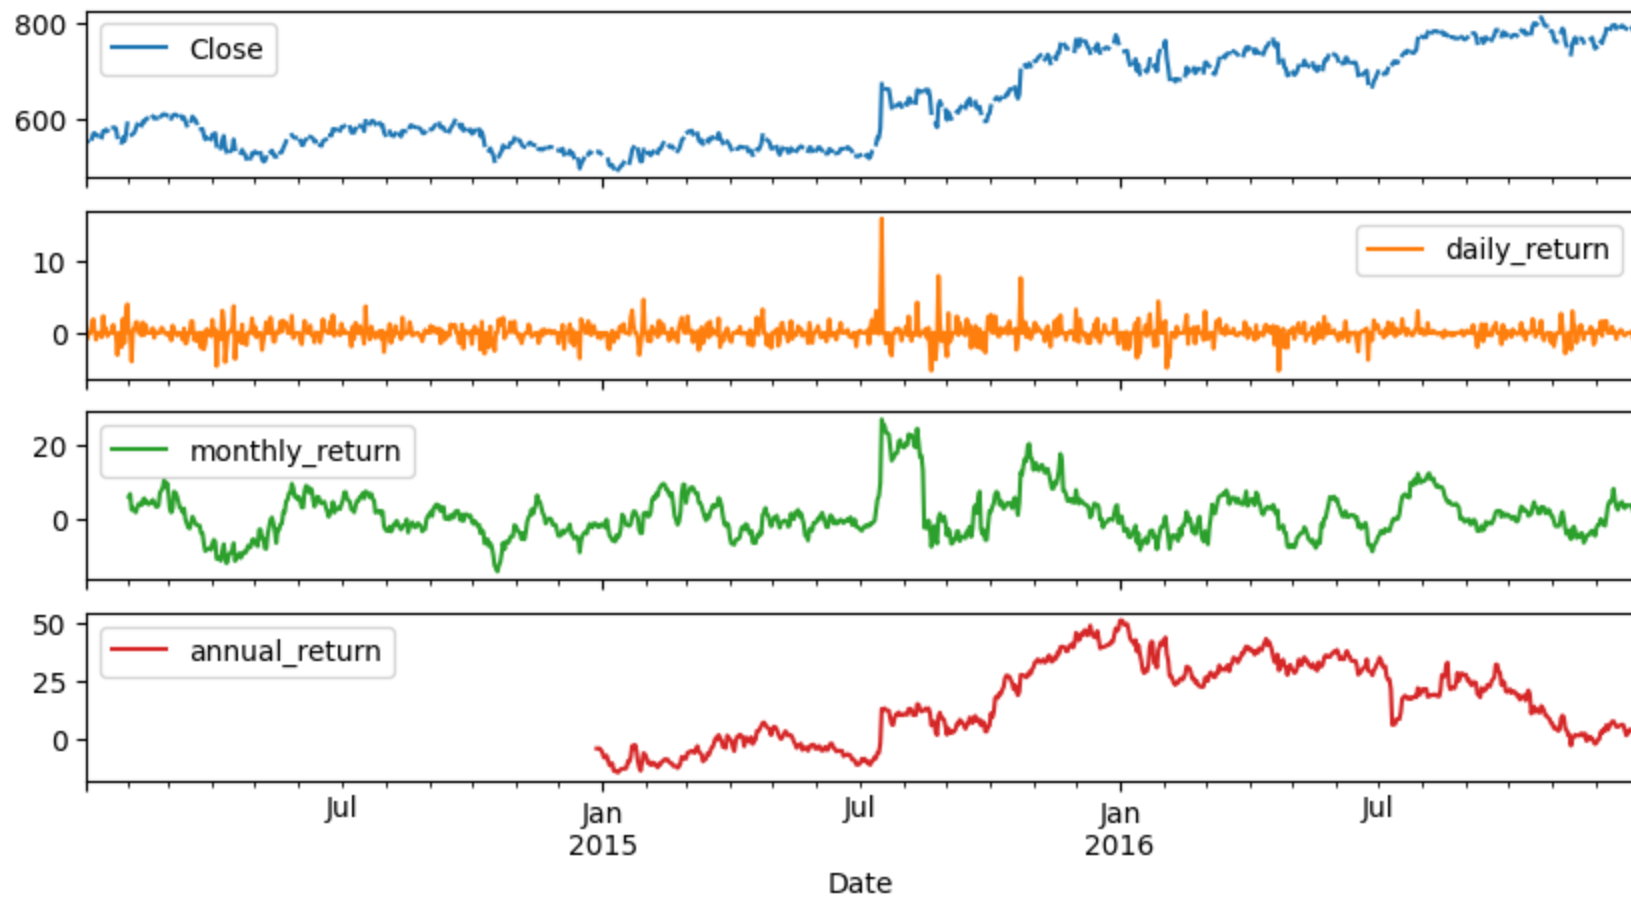
```
google['daily_return'] = google['Close'].pct_change(periods=1) * 100
```

C:\Users\yeiso\AppData\Local\Temp\ipykernel_19956\207313826.py:4: FutureWarning: The default fill_method='pad' in Series.pct_change is deprecated and will be removed in a future version. Call ffill before calling pct_change to retain current behavior and silence this warning.

```
google['monthly_return'] = google['Close'].pct_change(periods=30) * 100
```

C:\Users\yeiso\AppData\Local\Temp\ipykernel_19956\207313826.py:7: FutureWarning: The default fill_method='pad' in Series.pct_change is deprecated and will be removed in a future version. Call ffill before calling pct_change to retain current behavior and silence this warning.

```
google['annual_return'] = google['Close'].pct_change(periods=360) * 100
```



Chapter 2 - Basic Time Series Metrics & Resampling

This chapter dives deeper into the essential time series functionality made available through the pandas `DateTimeIndex`. It introduces resampling and how to compare different time series by normalizing their start points.

Compare the performance of several asset classes

```
In [ ]: # Import data here
# Import 'asset_classes.csv', using .read_csv() to parse dates in the 'DATE'
# column and set this column as the index, then assign the result to prices.
path = 'C:/Users/yeiso/OneDrive - Douglas College/0. DOUGLAS COLLEGE/3. Fund Machine Learning/0. Python Course DataCamp/Course-fun
prices.head()
```

Out[]:

	2013	2014	2015
0	20.08	NaN	NaN
1	19.78	39.59	50.17
2	19.86	40.12	49.13
3	19.40	39.93	49.21
4	19.66	40.92	48.59

In []:

```
# Import data here
prices = pd.read_csv(path+'asset_classes.csv', parse_dates=['DATE'], index_col='DATE')
```


```
# Inspect prices here
print(prices.info())
```

```
# Select first prices
first_prices = prices.iloc[0]
```

```
# Create normalized
normalized = prices.div(first_prices).mul(100)
```

```
# Plot normalized
normalized.plot()
plt.show()
```

```
<class 'pandas.core.frame.DataFrame'>
DatetimeIndex: 2469 entries, 2007-06-29 to 2017-06-26
Data columns (total 4 columns):
#   Column  Non-Null Count  Dtype
---  -
0    SP500    2469 non-null    float64
1    Bonds    2469 non-null    float64
2    Gold      2469 non-null    float64
3    Oil       2469 non-null    float64
dtypes: float64(4)
memory usage: 96.4 KB
None
```



Comparing stock prices with a benchmark

Compare the performance of various stocks against a benchmark. Learn more about the stock market by comparing the three largest stocks on the NYSE to the Dow Jones Industrial Average, which contains the 30 largest US companies.

The three largest companies on the NYSE are:

Company (Stock Ticker):

Johnson & Johnson (JNJ) Exxon Mobil (XOM) JP Morgan Chase (JPM)

```
In [ ]: # Import stock prices and index here
stocks = pd.read_csv(path+'nyse.csv', parse_dates=['date'], index_col='date')
dow_jones = pd.read_csv(path+'dow_jones.csv', parse_dates=['date'], index_col='date')
```

```
In [ ]: stocks.head()
```

Out[]:

JNJ JPM XOM

date			
2010-01-04	64.68	42.85	69.15
2010-01-05	63.93	43.68	69.42
2010-01-06	64.45	43.92	70.02
2010-01-07	63.99	44.79	69.80
2010-01-08	64.21	44.68	69.52

In []: `dow_jones.head()`

Out[]:

DJIA

date	
2010-01-04	10583.96
2010-01-05	10572.02
2010-01-06	10573.68
2010-01-07	10606.86
2010-01-08	10618.19

In []: *# Concatenate data and inspect result here*
Use pd.concat() along axis=1 to combine stocks and dow_jones
and assign the result to data. Inspect the .info() of data.
`data = pd.concat([stocks, dow_jones], axis=1)`
`print(data.info())`
`data.head()`

```
<class 'pandas.core.frame.DataFrame'>
DatetimeIndex: 1762 entries, 2010-01-04 to 2016-12-30
Data columns (total 4 columns):
#   Column  Non-Null Count  Dtype  
---  -
0    JNJ      1762 non-null    float64
1    JPM      1762 non-null    float64
2    XOM      1762 non-null    float64
3    DJIA     1762 non-null    float64
dtypes: float64(4)
memory usage: 68.8 KB
None
```


Out[]:

	JNJ	JPM	XOM	DJIA
date				
2010-01-04	64.68	42.85	69.15	10583.96
2010-01-05	63.93	43.68	69.42	10572.02
2010-01-06	64.45	43.92	70.02	10573.68
2010-01-07	63.99	44.79	69.80	10606.86
2010-01-08	64.21	44.68	69.52	10618.19

```
In [ ]: data.plot()  
plt.show();
```



```
In [ ]: # Normalize and plot your data here  
data.div(data.iloc[0]).mul(100).plot()  
plt.show();
```



Plot performance difference vs benchmark index

```
In [ ]: # Import stock data here
stocks = pd.read_csv('C:\\Users\\yeiso\\OneDrive - Douglas College\\0. DOUGLAS COLLEGE\\3. Fund Machine Learning\\0. Python Course\\stocks.csv')
stocks.head()
```

Out []: AAPL MSFT

date		
2007-01-03	11.97	29.86
2007-01-04	12.24	29.81
2007-01-05	12.15	29.64
2007-01-08	12.21	29.93
2007-01-09	13.22	29.96

```
In [ ]: # Import index here
sp500 = pd.read_csv(path+'sp500.csv', parse_dates=['date'], index_col='date')
sp500.head()
```


Out []:


	SP500
date	
2007-06-29	1503.35
2007-07-02	1519.43
2007-07-03	1524.87
2007-07-05	1525.40
2007-07-06	1530.44

```
In [ ]: # Concatenate stocks and index here
# Use pd.concat() to concatenate stocks and sp500 along axis=1,
# apply .dropna() to drop all missing values, and assign the result to data.
data = pd.concat([stocks, sp500], axis=1).dropna()
data.head()
```

Out []:

	AAPL	MSFT	SP500
date			
2007-06-29	17.43	29.47	1503.35
2007-07-02	17.32	29.74	1519.43
2007-07-03	18.17	30.02	1524.87
2007-07-05	18.96	29.99	1525.40
2007-07-06	18.90	29.97	1530.44

```
In [ ]: data.plot()
plt.show();
```




```
In [ ]: # Normalize data
normalized = data.div(data.iloc[0]).mul(100)
normalized.head()
```

```
Out[ ]:
```


	AAPL	MSFT	SP500
date			
2007-06-29	100.000000	100.000000	100.000000
2007-07-02	99.368904	100.916186	101.069611
2007-07-03	104.245554	101.866305	101.431470
2007-07-05	108.777969	101.764506	101.466724
2007-07-06	108.433735	101.696641	101.801976

```
In [ ]: normalized.plot()
plt.show();
```



```
In [ ]: # Create tickers
tickers = ['MSFT', 'AAPL']

# Subtract the normalized index from the normalized stock prices, and plot the result
# Select tickers from normalized, and subtract normalized['SP500']
# with keyword axis=0 to align the indexes, then plot the result.
normalized[tickers].sub(normalized['SP500'], axis=0).plot()
plt.show();
```



Now you can compare these stocks to the overall market so you can more easily spot trends and outliers.

Convert monthly to weekly data

```
In [ ]: import pandas as pd
```

```
# Set start and end dates
start = '2016-1-1'
end = '2016-2-29'
```

```
In [ ]: # Create monthly_dates here
# Create monthly_dates using pd.date_range with start,
# end and frequency alias 'M'.
monthly_dates = pd.date_range(start=start, end=end, freq='M')
print(monthly_dates)
```

```
DatetimeIndex(['2016-01-31', '2016-02-29'], dtype='datetime64[ns]', freq='M')
```

```
In [ ]: # Create and print monthly here
# Create and print the pd.Series monthly, passing the list [1, 2]
# as the data argument, and using monthly_dates as index.
```

```
monthly = pd.Series(data=[1, 2], index=monthly_dates)
print(monthly)
```

```
2016-01-31    1
2016-02-29    2
Freq: M, dtype: int64
```

```
In [ ]: # Create weekly_dates here
weekly_dates = pd.date_range(start=start, end=end, freq='W')
print(weekly_dates)
```

```
DatetimeIndex(['2016-01-03', '2016-01-10', '2016-01-17', '2016-01-24',
               '2016-01-31', '2016-02-07', '2016-02-14', '2016-02-21',
               '2016-02-28'],
              dtype='datetime64[ns]', freq='W-SUN')
```

```
In [ ]: # Print monthly, reindexed using weekly_dates
# Apply .reindex() to monthly three times: first without additional options,
# then with ffill and then with bfill, print()-ing each result.
print(monthly.reindex(weekly_dates))
```

```
2016-01-03    NaN
2016-01-10    NaN
2016-01-17    NaN
2016-01-24    NaN
2016-01-31    1.0
2016-02-07    NaN
2016-02-14    NaN
2016-02-21    NaN
2016-02-28    NaN
Freq: W-SUN, dtype: float64
```

```
In [ ]: print(monthly.reindex(weekly_dates, method='bfill'))
```

```
2016-01-03    1
2016-01-10    1
2016-01-17    1
2016-01-24    1
2016-01-31    1
2016-02-07    2
2016-02-14    2
2016-02-21    2
2016-02-28    2
Freq: W-SUN, dtype: int64
```

```
In [ ]: print(monthly.reindex(weekly_dates, method='ffill'))
```

```

2016-01-03      NaN
2016-01-10      NaN
2016-01-17      NaN
2016-01-24      NaN
2016-01-31      1.0
2016-02-07      1.0
2016-02-14      1.0
2016-02-21      1.0
2016-02-28      1.0
Freq: W-SUN, dtype: float64

```

Create weekly from monthly unemployment data

```

In [ ]: # Import data here
df = pd.read_csv('C:\\Users\\yeiso\\OneDrive - Douglas College\\0. DOUGLAS COLLEGE\\3. Fund Machine Learning\\0. Python Course Data\\unemployment.csv')
#df.drop(['Debt/GDP'], axis=1, inplace=True)
df.head()

```

```

Out[ ]:      Debt/GDP  Unemployment

```

date		
2010-01-01	87.00386	9.8
2010-02-01	NaN	9.8
2010-03-01	NaN	9.9
2010-04-01	88.67047	9.9
2010-05-01	NaN	9.6

all example complete!!!

```

In [ ]: # Import data here
df = pd.read_csv('C:\\Users\\yeiso\\OneDrive - Douglas College\\0. DOUGLAS COLLEGE\\3. Fund Machine Learning\\0. Python Course Data\\unemployment.csv')
data=df

# Show first five rows of weekly series
print(data.asfreq('W').head(5))

# Show first five rows of weekly series with bfill option
print(data.asfreq('W', method='bfill').head(5))

# Create weekly series with ffill option and show first five rows
weekly_ffill = data.asfreq('W', method='ffill')
print(weekly_ffill.head())

# Plot weekly_fill starting 2015 here


```

```
weekly_ffill['2015:'].plot()  
plt.show();
```

	Debt/GDP	Unemployment
date		
2010-01-03	NaN	NaN
2010-01-10	NaN	NaN
2010-01-17	NaN	NaN
2010-01-24	NaN	NaN
2010-01-31	NaN	NaN

	Debt/GDP	Unemployment
date		
2010-01-03	NaN	9.8
2010-01-10	NaN	9.8
2010-01-17	NaN	9.8
2010-01-24	NaN	9.8
2010-01-31	NaN	9.8

	Debt/GDP	Unemployment
date		
2010-01-03	87.00386	9.8
2010-01-10	87.00386	9.8
2010-01-17	87.00386	9.8
2010-01-24	87.00386	9.8
2010-01-31	87.00386	9.8



Use interpolation to create weekly employment data

```
In [ ]: import pandas as pd
import matplotlib.pyplot as plt

# Import data here
df = pd.read_csv('C:\\Users\\yeiso\\OneDrive - Douglas College\\0. DOUGLAS COLLEGE\\3. Fund Machine Learning\\0. Python Course Data\\data.csv')
#df.drop(['Debt/GDP'], axis=1, inplace=True)
df.head()
```



Out []: **Debt/GDP Unemployment**

date		
2010-01-01	87.00386	9.8
2010-02-01	NaN	9.8
2010-03-01	NaN	9.9
2010-04-01	88.67047	9.9
2010-05-01	NaN	9.6

```
In [ ]: monthly = df
# Inspect data here
print(monthly.info())

<class 'pandas.core.frame.DataFrame'>
DatetimeIndex: 89 entries, 2010-01-01 to 2017-05-01
Data columns (total 2 columns):
#   Column          Non-Null Count  Dtype
---  -
0   Debt/GDP         29 non-null    float64
1   Unemployment     89 non-null    float64
dtypes: float64(2)
memory usage: 2.1 KB
None
```

```
In [ ]: monthly.plot()
plt.show();
```



```
In [ ]: # Create weekly dates
# Create a pd.date_range() with weekly dates, using the .min() and .max()
# of the index of monthly as start and end, respectively,
# and assign the result to weekly_dates.
weekly_dates = pd.date_range(start=monthly.index.min(), end=monthly.index.max(), freq='W')
weekly_dates
```

```
Out[ ]: DatetimeIndex(['2010-01-03', '2010-01-10', '2010-01-17', '2010-01-24',
                     '2010-01-31', '2010-02-07', '2010-02-14', '2010-02-21',
                     '2010-02-28', '2010-03-07',
                     ...,
                     '2017-02-26', '2017-03-05', '2017-03-12', '2017-03-19',
                     '2017-03-26', '2017-04-02', '2017-04-09', '2017-04-16',
                     '2017-04-23', '2017-04-30'],
                    dtype='datetime64[ns]', length=383, freq='W-SUN')
```


```
In [ ]: # Reindex monthly to weekly data
# Apply .reindex() using weekly_dates to monthly
# and assign the output to weekly.
weekly = monthly.reindex(weekly_dates)
weekly.tail()
```

Out []:

	Debt/GDP	Unemployment
--	----------	--------------

2017-04-02	NaN	NaN
2017-04-09	NaN	NaN
2017-04-16	NaN	NaN
2017-04-23	NaN	NaN
2017-04-30	NaN	NaN

```
In [ ]: weekly.plot()  
plt.show();
```



```
In [ ]: # Create ffill and interpolated columns  
weekly['ffill'] = weekly.UNRATE.ffill()  
weekly['interpolated'] = weekly.UNRATE.interpolate()
```


```
In [ ]: weekly.tail()
```

Out[]:

	Debt/GDP	Unemployment
--	----------	--------------

2017-04-02	NaN	NaN
2017-04-09	NaN	NaN
2017-04-16	NaN	NaN
2017-04-23	NaN	NaN
2017-04-30	NaN	NaN

```
In [ ]: # Plot weekly
weekly.plot()
plt.show();
```



Interpolate debt/GDP and compare to unemployment

Since you have learned how to interpolate time series, you can now apply this new skill to the quarterly debt/GDP series, and compare the result to the monthly unemployment rate.

```
In [ ]: # Import & inspect data here
df = pd.read_csv('C:\\Users\\yeiso\\OneDrive - Douglas College\\0. DOUGLAS COLLEGE\\3. Fund Machine Learning\\0. Python Course Data\\data.csv')
print(df.info())
df.head()
```

```
<class 'pandas.core.frame.DataFrame'>
DatetimeIndex: 89 entries, 2010-01-01 to 2017-05-01
Data columns (total 2 columns):
#   Column          Non-Null Count  Dtype
---  -
0   Debt/GDP         29 non-null    float64
1   Unemployment      89 non-null    float64
dtypes: float64(2)
memory usage: 4.1 KB
None
```

```
Out [ ]:      Debt/GDP  Unemployment
```

date		
2010-01-01	87.00386	9.8
2010-02-01	NaN	9.8
2010-03-01	NaN	9.9
2010-04-01	88.67047	9.9
2010-05-01	NaN	9.6

```
In [ ]: # Interpolate and inspect here
interpolated = df.interpolate()
print(interpolated.info())
interpolated.head()
```


```
<class 'pandas.core.frame.DataFrame'>
DatetimeIndex: 89 entries, 2010-01-01 to 2017-05-01
Data columns (total 2 columns):
#   Column          Non-Null Count  Dtype
---  -
0   Debt/GDP         89 non-null    float64
1   Unemployment      89 non-null    float64
dtypes: float64(2)
memory usage: 4.1 KB
None
```

Out[]:

Debt/GDP Unemployment

date		
2010-01-01	87.003860	9.8
2010-02-01	87.559397	9.8
2010-03-01	88.114933	9.9
2010-04-01	88.670470	9.9
2010-05-01	89.135103	9.6

```
In [ ]: # Plot interpolated data here
interpolated.plot(secondary_y='Unemployment')
plt.show();
```



Downsampling & aggregation

Compare weekly, monthly and annual ozone trends for NYC-LA

Downsample and aggregate time series ozone data for both NYC and LA since 2000 to compare the air quality trend at weekly, monthly and annual frequencies and explore how different resampling periods impact the visualization.

```
In [ ]: # Import and inspect data here
ozone = pd.read_csv('C:\\Users\\yeiso\\OneDrive - Douglas College\\0. DOUGLAS COLLEGE\\3. Fund Machine Learning\\0. Python Course
print(ozone.info())
ozone.head()
```

```
<class 'pandas.core.frame.DataFrame'>
DatetimeIndex: 6291 entries, 2000-01-01 to 2017-03-31
Data columns (total 2 columns):
#   Column      Non-Null Count  Dtype
---  -
0   Los Angeles  5488 non-null   float64
1   New York     6167 non-null   float64
dtypes: float64(2)
memory usage: 147.4 KB
None
```


```
Out[ ]:           Los Angeles  New York
```

date		
2000-01-01	0.008375	0.004032
2000-01-02	NaN	0.009486
2000-01-03	NaN	0.005580
2000-01-04	0.005500	0.008717
2000-01-05	0.005000	0.013754

```
In [ ]: # Import and inspect data here
stocks = pd.read_csv('C:\\Users\\yeiso\\OneDrive - Douglas College\\0. DOUGLAS COLLEGE\\3. Fund Machine Learning\\0. Python Course
print(stocks.info())
```

```
<class 'pandas.core.frame.DataFrame'>
DatetimeIndex: 504 entries, 2015-01-02 to 2016-12-30
Data columns (total 2 columns):
#   Column  Non-Null Count  Dtype
---  -
0   FB       504 non-null    float64
1   GOOG     504 non-null    float64
dtypes: float64(2)
memory usage: 11.8 KB
None
```

```
In [ ]: # Calculate and plot the monthly averages
monthly_average = stocks.resample('M').mean()
monthly_average.plot(subplots=True)
plt.show();
```



Compare quarterly GDP growth rate and stock returns

With your new skill to downsample and aggregate time series, you can compare higher-frequency stock price series to lower-frequency economic time series.

As a first example, let's compare the quarterly GDP growth rate to the quarterly rate of return on the (resampled) Dow Jones Industrial index of 30 large US stocks.

GDP growth is reported at the beginning of each quarter for the previous quarter. To calculate matching stock returns, you'll resample the stock index to quarter start frequency using the alias 'QS', and aggregating using the `.first()` observations.


```

In [ ]: # Import and inspect gdp_growth here
gdp_growth = pd.read_csv('C:\\Users\\yeiso\\OneDrive - Douglas College\\0. DOUGLAS COLLEGE\\3. Fund Machine Learning\\0. Python Co
gdp_growth.info()

<class 'pandas.core.frame.DataFrame'>
DatetimeIndex: 41 entries, 2007-01-01 to 2017-01-01
Data columns (total 1 columns):
#   Column      Non-Null Count  Dtype
---  ---
0    gdp_growth  41 non-null      float64
dtypes: float64(1)
memory usage: 656.0 bytes

In [ ]: # Import and inspect djia here
djia = pd.read_csv('C:\\Users\\yeiso\\OneDrive - Douglas College\\0. DOUGLAS COLLEGE\\3. Fund Machine Learning\\0. Python Course D
djia.info()

<class 'pandas.core.frame.DataFrame'>
DatetimeIndex: 2610 entries, 2007-06-29 to 2017-06-29
Data columns (total 1 columns):
#   Column      Non-Null Count  Dtype
---  ---
0    djia        2519 non-null    float64
dtypes: float64(1)
memory usage: 40.8 KB

In [ ]: # Calculate djia quarterly returns here
djia_quarterly = djia.resample('QS').first()
djia_quarterly_return = djia_quarterly.pct_change().mul(100)

In [ ]: # Concatenate, rename and plot djia_quarterly_return and gdp_growth here
# Use pd.concat() to concatenate gdp_growth and djia_quarterly_return
# along axis=1, and assign to data. Rename the columns using .columns
# and the new labels 'gdp' and 'djia', then .plot() the results.
data = pd.concat([gdp_growth, djia_quarterly_return], axis=1)
data.columns = ['gdp', 'djia']

In [ ]: data.plot()
plt.show();

```



Visualize monthly mean, median and standard deviation of S&P500 returns

You have also learned how to calculate several aggregate statistics from upsampled data.

Let's use this to explore how the monthly mean, median and standard deviation of daily S&P500 returns have trended over the last 10 years.

```
In [ ]: # Import data here
sp500 = pd.read_csv('C:\\Users\\yeiso\\OneDrive - Douglas College\\0. DOUGLAS COLLEGE\\3. Fund Machine Learning\\0. Python Course
print(sp500.info())
sp500.head()

<class 'pandas.core.frame.DataFrame'>
DatetimeIndex: 2395 entries, 2007-06-29 to 2016-12-30
Data columns (total 1 columns):
#   Column  Non-Null Count  Dtype
---  -
0    SP500    2395 non-null    float64
dtypes: float64(1)
memory usage: 37.4 KB
None
```

Out[]: **SP500**

date	
2007-06-29	1503.35
2007-07-02	1519.43
2007-07-03	1524.87
2007-07-05	1525.40
2007-07-06	1530.44

```
In [ ]: # Calculate daily returns here
# Convert sp500 to a pd.Series() using .squeeze(),
# and apply .pct_change() to calculate daily_returns.
daily_returns = sp500.squeeze().pct_change()
daily_returns
```

```
Out[ ]: date
2007-06-29      NaN
2007-07-02      0.010696
2007-07-03      0.003580
2007-07-05      0.000348
2007-07-06      0.003304
...
2016-12-23      0.001252
2016-12-27      0.002248
2016-12-28     -0.008357
2016-12-29     -0.000293
2016-12-30     -0.004637
Name: SP500, Length: 2395, dtype: float64
```


```
In [ ]: # Resample and calculate statistics
# .resample() daily_returns to month-end frequency (alias: 'M'),
# and apply .agg() to calculate 'mean', 'median', and 'std'.
# Assign the result to stats.
stats = daily_returns.resample('M').agg(['mean', 'median', 'std'])
stats
```

Out[]: mean median std

date			
2007-06-30	NaN	NaN	NaN
2007-07-31	-0.001490	0.000921	0.010908
2007-08-31	0.000668	0.001086	0.015261
2007-09-30	0.001900	0.000202	0.010000
2007-10-31	0.000676	-0.000265	0.008719
...
2016-08-31	-0.000047	-0.000796	0.003562
2016-09-30	-0.000019	-0.000019	0.009146
2016-10-31	-0.000925	-0.001376	0.004160
2016-11-30	0.001623	0.000808	0.006675
2016-12-31	0.000871	0.001252	0.005040

115 rows × 3 columns

```
In [ ]: # Plot stats here
stats.plot()
plt.show();
```



In []:

Chapter 3 - Window Functions: Rolling & Expanding Metrics

Rolling average air quality since 2010 for new york city

To practice rolling window functions, you'll start with air quality trends for New York City since 2010. In particular, you'll be using the daily Ozone concentration levels provided by the Environmental Protection Agency to calculate & plot the 90 and 360 day rolling average.

```
In [ ]: import pandas as pd
import matplotlib.pyplot as plt
```

```
In [ ]: # Import and inspect ozone data here
data = pd.read_csv('C:\\Users\\yeiso\\OneDrive - Douglas College\\0. DOUGLAS COLLEGE\\3. Fund Machine Learning\\0. Python Course D
print(data.info())
data.head()
```


```
<class 'pandas.core.frame.DataFrame'>
DatetimeIndex: 6291 entries, 2000-01-01 to 2017-03-31
Data columns (total 1 columns):
#   Column  Non-Null Count  Dtype
---  -
0   Ozone    6167 non-null      float64
dtypes: float64(1)
memory usage: 98.3 KB
None
```

Ozone	
date	
2000-01-01	0.004032
2000-01-02	0.009486
2000-01-03	0.005580
2000-01-04	0.008717
2000-01-05	0.013754

```
In [ ]: # Calculate 90d and 360d rolling mean for the last price
data['90D'] = data.Ozone.rolling(window='90D').mean()
data['360D'] = data.Ozone.rolling(window='360D').mean()
data.head()
```

	Ozone	90D	360D
date			
2000-01-01	0.004032	0.004032	0.004032
2000-01-02	0.009486	0.006759	0.006759
2000-01-03	0.005580	0.006366	0.006366
2000-01-04	0.008717	0.006954	0.006954
2000-01-05	0.013754	0.008314	0.008314

```
In [ ]: # Plot data
data['2010:'].plot()
plt.title('New York City')
plt.show();
```



Rolling 360-day median & std. deviation for nyc ozone data since 2000

Calculate several rolling statistics using the `.agg()` method, similar to `.groupby()`.

Let's take a closer look at the air quality history of NYC using the Ozone data you have seen before. The daily data are very volatile, so using a longer term rolling average can help reveal a longer term trend.

You'll be using a 360 day rolling window, and `.agg()` to calculate the rolling median and standard deviation for the daily average ozone values since 2000.

```
In [ ]: # Import and inspect ozone data here
data = pd.read_csv('C:\\Users\\yeiso\\OneDrive - Douglas College\\0. DOUGLAS COLLEGE\\3. Fund Machine Learning\\0. Python Course D
data.head()
```

Ozone	
date	
2000-01-01	0.004032
2000-01-02	0.009486
2000-01-03	0.005580
2000-01-04	0.008717
2000-01-05	0.013754

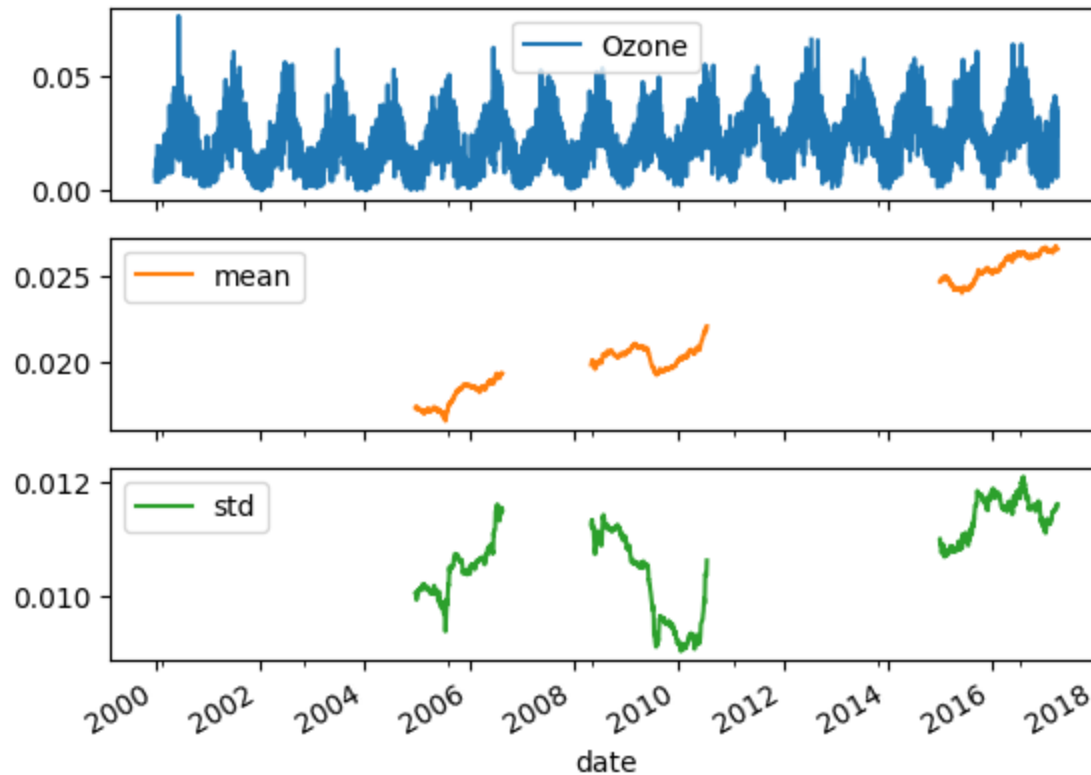
```
In [ ]: # Calculate the rolling mean and std here
rolling_stats = data.Ozone.rolling(360).agg(['mean', 'std'])
rolling_stats.tail()
```

	mean	std
date		
2017-03-27	0.026629	0.011599
2017-03-28	0.026583	0.011617
2017-03-29	0.026584	0.011617
2017-03-30	0.026599	0.011613
2017-03-31	0.026607	0.011618

```
In [ ]: # Join rolling_stats with ozone data
stats = data.join(rolling_stats)
stats.head()
```

	Ozone	mean	std
date			
2000-01-01	0.004032	NaN	NaN
2000-01-02	0.009486	NaN	NaN
2000-01-03	0.005580	NaN	NaN
2000-01-04	0.008717	NaN	NaN
2000-01-05	0.013754	NaN	NaN


```
In [ ]: # Plot stats
stats.plot(subplots=True)
plt.show();
```



Rolling quantiles for daily air quality in nyc

Calculate rolling quantiles to describe changes in the dispersion of a time series over time in a way that is less sensitive to outliers than using the mean and standard deviation.

Let's calculate rolling quantiles - at 10%, 50% (median) and 90% - of the distribution of daily average ozone concentration in NYC using a 360-day rolling window.

```
In [ ]: # Import and inspect ozone data here
data = pd.read_csv('C:\\Users\\yeiso\\OneDrive - Douglas College\\0. DOUGLAS COLLEGE\\3. Fund Machine Learning\\0. Python Course D
data.head()
```

Ozone

date

2000-01-01	0.004032
2000-01-02	0.009486
2000-01-03	0.005580
2000-01-04	0.008717
2000-01-05	0.013754


```
In [ ]: # Resample, interpolate and inspect ozone data here
data = data.resample('D').interpolate()
print(data.info())

<class 'pandas.core.frame.DataFrame'>
DatetimeIndex: 6300 entries, 2000-01-01 to 2017-03-31
Freq: D
Data columns (total 1 columns):
#   Column  Non-Null Count  Dtype
---  -
0   Ozone    6300 non-null    float64
dtypes: float64(1)
memory usage: 98.4 KB
None
```

```
In [ ]: # Create the rolling window
rolling = data.Ozone.rolling(360)

# Insert the rolling quantiles to the monthly returns
data['q10'] = rolling.quantile(0.1)
data['q50'] = rolling.quantile(0.5)
data['q90'] = rolling.quantile(0.9)

# Plot the data
data.plot()
plt.show();
```



Expanding window functions with pandas

Cumulative sum vs .diff()

Expanding windows allow you to run cumulative calculations.

The cumulative sum method has in fact the opposite effect of the `.diff()` method that you came across earlier. To illustrate this, let's use the Google stock price time series, create the differences between prices, and reconstruct the series using the cumulative sum.

```
In [ ]: import pandas as pd
import matplotlib.pyplot as plt

# Import data here
df = pd.read_csv('C:\\Users\\yeiso\\OneDrive - Douglas College\\0. DOUGLAS COLLEGE\\3. Fund Machine Learning\\0. Python Course Data\\Google Stock Price.csv')
df.head()
```

Close

Date

Date	Close
2014-01-02	556.00
2014-01-03	551.95
2014-01-06	558.10
2014-01-07	568.86
2014-01-08	570.04

```
In [ ]: # Calculate differences
differences = df.diff().dropna()
differences.head()
```

Close

Date

Date	Close
2014-01-03	-4.05
2014-01-06	6.15
2014-01-07	10.76
2014-01-08	1.18
2014-01-09	-5.49

```
In [ ]: # Select start price
start_price = df.first('D')
start_price
```

C:\Users\yeiso\AppData\Local\Temp\ipykernel_38792\863676461.py:2: FutureWarning: first is deprecated and will be removed in a future version. Please create a mask and filter using `.loc` instead

```
start_price = df.first('D')
```

Close

Date

Date	Close
2014-01-02	556.0

```
In [ ]: #este era el codigo original pero no funciona
#cumulative_sum = start_price.append(differences).cumsum()
#cumulative_sum.head()
```

```
#en vez de eso... use este!
# Calculate cumulative sum
cumulative_sum = pd.concat([start_price, differences]).cumsum()
cumulative_sum.head()
```

	Close
Date	
2014-01-02	556.00
2014-01-03	551.95
2014-01-06	558.10
2014-01-07	568.86
2014-01-08	570.04

```
In [ ]: # Validate cumulative sum equals data
print(df.equals(cumulative_sum))
```

True

Cumulative return on \$1,000 invested in google vs apple I

To put your new ability to do cumulative return calculations to practical use, let's compare how much \$1,000 would be worth if invested in Google ('GOOG') or Apple ('AAPL') in 2010.

```
In [ ]: # Import data here
df = pd.read_csv('C:\\Users\\yeiso\\OneDrive - Douglas College\\0. DOUGLAS COLLEGE\\3. Fund Machine Learning\\0. Python Course Data\\AAPL.csv')
df.head()
```

	AAPL	GOOG
Date		
2010-01-04	NaN	313.06
2010-01-05	NaN	311.68
2010-01-06	NaN	303.83
2010-01-07	NaN	296.75
2010-01-08	NaN	300.71

```
In [ ]: # Define your investment
investment = 1000
```


```
In [ ]: # Calculate the daily returns here
returns = df.pct_change()
returns
```

	AAPL	GOOG
Date		
2010-01-04	NaN	NaN
2010-01-05	NaN	-0.004408
2010-01-06	NaN	-0.025186
2010-01-07	NaN	-0.023303
2010-01-08	NaN	0.013345
...
2017-05-24	-0.002991	0.006471
2017-05-25	0.003456	0.015268
2017-05-26	-0.001690	0.001991
2017-05-30	0.000391	0.004540
2017-05-31	-0.005922	-0.011292

1864 rows × 2 columns

```
In [ ]: # Calculate the cumulative returns here
returns_plus_one = returns + 1
cumulative_return = returns_plus_one.cumprod()
```

```
In [ ]: # Calculate and plot the investment return here
cumulative_return.mul(investment).plot()
plt.show();
```



Cumulative return on \$1,000 invested in google vs apple II

Apple outperformed Google over the entire period, but this may have been different over various 1-year sub periods, so that switching between the two stocks might have yielded an even better result.

To analyze this, calculate that cumulative return for rolling 1-year periods, and then plot the returns to see when each stock was superior.

```
In [ ]: # Import numpy
import numpy as np
```

```
In [ ]: # Define a multi_period_return function
def multi_period_return(period_returns):
    return np.prod(period_returns + 1) - 1
```

```
In [ ]: # Import data here
df = pd.read_csv('C:\\Users\\yeiso\\OneDrive - Douglas College\\0. DOUGLAS COLLEGE\\3. Fund Machine Learning\\0. Python Course Data\\AAPL.csv')
df.head()
```

AAPL GOOG

Date		
2010-01-04	NaN	313.06
2010-01-05	NaN	311.68
2010-01-06	NaN	303.83
2010-01-07	NaN	296.75
2010-01-08	NaN	300.71

```
In [ ]: # Calculate daily returns
daily_returns = df.pct_change()
daily_returns.head()
```

AAPL GOOG


Date		
2010-01-04	NaN	NaN
2010-01-05	NaN	-0.004408
2010-01-06	NaN	-0.025186
2010-01-07	NaN	-0.023303
2010-01-08	NaN	0.013345

```
In [ ]: # Calculate rolling annual returns
rolling_annual_returns = daily_returns.rolling('360D').apply(multi_period_return)
rolling_annual_returns.tail()
```

AAPL GOOG

Date		
2017-05-24	0.528052	0.303415
2017-05-25	0.533333	0.323315
2017-05-26	0.538254	0.320434
2017-05-30	0.569342	0.350998
2017-05-31	0.560049	0.335742


```
In [ ]: # Plot rolling_annual_returns
rolling_annual_returns.mul(100).plot()
plt.show();
```



Case study: S&P500 price simulation

In this exercise, you'll build your own random walk by drawing random numbers from the normal distribution with the help of numpy.

Random walk I You'll build your own random walk by drawing random numbers from the normal distribution with the help of numpy.

```
In [ ]: from numpy.random import normal, seed
from scipy.stats import norm
```

```
In [ ]: # Set seed here
seed = 42
np.random.seed(seed)

# Create random_walk
# Use normal to generate 2,500 random returns with the parameters
# loc=.001, scale=.01 and assign this to random_walk.
random_walk = normal(loc=.001, scale=0.01, size=2500)
```


```

# Convert random_walk to pd.Series
random_walk = pd.Series(random_walk)

# Create random_prices
# Create random_prices by adding 1 to random_walk
# and calculating the cumulative product.
random_prices = random_walk.add(1).cumprod()

# Plot random_prices here
# Multiply random_prices by 1,000
# and plot the result for a price series starting at 1,000.
random_prices.mul(1000).plot()
plt.show();

```



Random walk II

You'll build a random walk using historical returns from Facebook's stock price since IPO through the end of May 31, 2017. Then you'll simulate an alternative random price path in the next exercise.

```

In [ ]: # Import data here
df = pd.read_csv('C:\\Users\\yeiso\\OneDrive - Douglas College\\0. DOUGLAS COLLEGE\\3. Fund Machine Learning\\0. Python Course Dat
df.index.name = 'date'

```

```
df.columns = ['price']  
df.head()
```

	price
date	
2012-05-17	38.00
2012-05-18	38.23
2012-05-21	34.03
2012-05-22	31.00
2012-05-23	32.00

```
In [ ]: fb = df['price']  
print(type(fb))  
fb
```

```
<class 'pandas.core.series.Series'>  
date  
2012-05-17    38.00  
2012-05-18    38.23  
2012-05-21    34.03  
2012-05-22    31.00  
2012-05-23    32.00  
...  
2017-05-24    150.04  
2017-05-25    151.96  
2017-05-26    152.13  
2017-05-30    152.38  
2017-05-31    151.46  
Name: price, Length: 1267, dtype: float64
```

```
In [ ]: import pandas as pd  
from numpy.random import choice, random  
import matplotlib.pyplot as plt  
import seaborn as sns  
  
# Set seed here  
#sometimes using >>> seed(42) <<< could works!  
seed = 42  
np.random.seed(seed)
```

```
In [ ]: # Calculate daily_returns here  
daily_returns = fb.pct_change().dropna()  
print(daily_returns)
```

```
date
2012-05-18    0.006053
2012-05-21   -0.109861
2012-05-22   -0.089039
2012-05-23    0.032258
2012-05-24    0.032188
...
2017-05-24    0.013305
2017-05-25    0.012797
2017-05-26    0.001119
2017-05-30    0.001643
2017-05-31   -0.006038
Name: price, Length: 1266, dtype: float64
```

```
In [ ]: # Get n_obs
n_obs = daily_returns.count()
print(n_obs)
```

```
1266
```

```
In [ ]: # Create random_walk
random_walk = choice(daily_returns, size=n_obs)
random_walk
```

```
array([-0.00637783, -0.00854701,  0.00833254, ..., -0.00832266,
       -0.00044709, -0.00940827])
```

```
In [ ]: # Convert random_walk to pd.series
random_walk = pd.Series(random_walk)
```

```
In [ ]: # Plot random_walk distribution
sns.distplot(random_walk)
plt.show();
```


C:\Users\yeiso\AppData\Local\Temp\ipykernel_38792\3014660000.py:2: UserWarning:

`distplot` is a deprecated function and will be removed in seaborn v0.14.0.

Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).

For a guide to updating your code to use the new functions, please see <https://gist.github.com/mwaskom/de44147ed2974457ad6372750bbe5751>

```
sns.distplot(random_walk)
```



Random walk III

In this exercise, you'll complete your random walk simulation using Facebook stock returns over the last five years. You'll start off with a random sample of returns like the one you've generated during the last exercise and use it to create a random stock price path.

```
In [ ]: # Select fb start price here
start = fb.price.first('D')

# Add 1 to random walk and append to start
random_walk = random_walk.add(1)
random_price = start.append(random_walk)

# Calculate cumulative product here
random_price = random_price.cumprod()

# Insert into fb and plot
fb['random'] = random_price

fb.plot()
plt.show()
```

Annual return correlations among several stocks

You have seen in the video how to calculate correlations, and visualize the result.

In this exercise, we have provided you with the historical stock prices for Apple (AAPL), Amazon (AMZN), IBM (IBM), WalMart (WMT), and Exxon Mobile (XOM) for the last 4,000 trading days from July 2001 until the end of May 2017.

You'll calculate the year-end returns, the pairwise correlations among all stocks, and visualize the result as an annotated heatmap.

```
In [ ]: # Import data here
df = pd.read_csv('C:\\Users\\yeiso\\OneDrive - Douglas College\\0. DOUGLAS COLLEGE\\3. Fund Machine Learning\\0. Python Course Data\\stock_prices.csv')
df.head()
```

	AAPL	AMZN	IBM	WMT	XOM
Date					
2001-07-05	1.66	15.27	NaN	NaN	NaN
2001-07-06	1.57	15.27	106.50	47.34	43.40
2001-07-09	1.62	15.81	104.72	48.25	43.36
2001-07-10	1.51	15.61	101.96	47.50	42.88
2001-07-11	1.61	15.34	103.85	48.85	42.48

```
In [ ]: data = df
# Inspect data here
print(data.info())

# Calculate year-end prices here
annual_prices = data.resample('A').last()


# Calculate annual returns here
annual_returns = annual_prices.pct_change()

# Calculate and print the correlation matrix here
correlations = annual_returns.corr()
print(correlations)

# Visualize the correlations as heatmap here
sns.heatmap(correlations, annot=True)
plt.show();
```

```
<class 'pandas.core.DataFrame'>
DatetimeIndex: 4001 entries, 2001-07-05 to 2017-05-31
Data columns (total 5 columns):
#   Column  Non-Null Count  Dtype
---  -
0   AAPL    4000 non-null      float64
1   AMZN    4000 non-null      float64
2   IBM     4000 non-null      float64
3   WMT     4000 non-null      float64
4   XOM     4000 non-null      float64
dtypes: float64(5)
memory usage: 187.5 KB
None
```

	AAPL	AMZN	IBM	WMT	XOM
AAPL	1.000000	0.208731	0.460568	-0.183553	0.336413
AMZN	0.208731	1.000000	0.346407	-0.367620	-0.133965
IBM	0.460568	0.346407	1.000000	0.155445	0.367253
WMT	-0.183553	-0.367620	0.155445	1.000000	0.178833
XOM	0.336413	-0.133965	0.367253	0.178833	1.000000



Chapter 4 - Putting it all together

```
In [ ]: #import requiered libraries
```

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
```

```
In [ ]: listings_nyse = pd.read_excel('C:\\Users\\yeiso\\OneDrive - Douglas College\\0. DOUGLAS COLLEGE\\3. Fund Machine Learning\\0. Pyth
listings_amex = pd.read_excel('C:\\Users\\yeiso\\OneDrive - Douglas College\\0. DOUGLAS COLLEGE\\3. Fund Machine Learning\\0. Pyth
listings_nasdaq = pd.read_excel('C:\\Users\\yeiso\\OneDrive - Douglas College\\0. DOUGLAS COLLEGE\\3. Fund Machine Learning\\0. Py
```

```
c:\Users\yeiso\AppData\Local\Programs\Python\Python312\Lib\site-packages\openpyxl\worksheet\_reader.py:329: UserWarning: Unknown
extension is not supported and will be removed
  warn(msg)
c:\Users\yeiso\AppData\Local\Programs\Python\Python312\Lib\site-packages\openpyxl\worksheet\_reader.py:329: UserWarning: Unknown
extension is not supported and will be removed
  warn(msg)
c:\Users\yeiso\AppData\Local\Programs\Python\Python312\Lib\site-packages\openpyxl\worksheet\_reader.py:329: UserWarning: Unknown
extension is not supported and will be removed
  warn(msg)
```

```
In [ ]: listings_nyse['Exchange'] = 'nyse'
listings_amex['Exchange'] = 'amex'
listings_nasdaq['Exchange'] = 'nasdaq'
```

```
In [ ]: listings = pd.concat([listings_amex, listings_nasdaq, listings_nyse], axis=0)
listings.reset_index(inplace=True)
listings.drop(['index'], axis=1, inplace=True)
listings['Market Capitalization'] /= 1e6
```

```
In [ ]: print(listings.info())

# Move 'stock symbol' into the index
listings.set_index('Stock Symbol', inplace=True)

# Drop rows with missing 'sector' data
listings.dropna(subset=['Sector'], inplace=True)

# Select companies with IPO Year befor 2019
listings = listings[listings['IPO Year'] < 2019]

# Inspect the new listings data
print(listings.info())

# Show the number of companies per sector
print(listings.groupby('Sector').size().sort_values(ascending=False))
```



```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 9441 entries, 0 to 9440
Data columns (total 8 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Stock Symbol          9441 non-null  object
1   Company Name          9441 non-null  object
2   Last Sale             9237 non-null  float64
3   Market Capitalization 9441 non-null  float64
4   IPO Year              4083 non-null  float64
5   Sector                6531 non-null  object
6   Industry              6531 non-null  object
7   Exchange              9441 non-null  object
```

dtypes: float64(3), object(5)

memory usage: 590.2+ KB

None

```
<class 'pandas.core.frame.DataFrame'>
```

Index: 2901 entries, WBAI to ZTO

Data columns (total 7 columns):

```
#   Column                Non-Null Count  Dtype
---  -
0   Company Name          2901 non-null  object
1   Last Sale             2901 non-null  float64
2   Market Capitalization 2901 non-null  float64
3   IPO Year              2901 non-null  float64
4   Sector                2901 non-null  object
5   Industry              2901 non-null  object
6   Exchange              2901 non-null  object
```

dtypes: float64(3), object(4)

memory usage: 181.3+ KB

None

Sector

```
Consumer Services    720
Finance              396
Energy               339
Technology            324
Public Utilities     228
Basic Industries     198
Capital Goods        183
Consumer Non-Durables 141
Health Care          129
Transportation        111
Miscellaneous         75
Consumer Durables     57
```

dtype: int64

Select and inspect index components

Now that you have imported and cleaned the listings data, you can proceed to select the index components as the largest company for each sector by market capitalization.

You'll also have the opportunity to take a closer look at the components, their last market value, and last price.

```
In [ ]: components = listings.groupby('Sector')['Market Capitalization'].nlargest(1)

# Print components, sorted by market cap
print(components.sort_values(ascending=False))

# Select stock symbols and print the result
tickers = components.index.get_level_values('Stock Symbol')
print(tickers)

# Print company name, market cap, and last price for each components
info_cols = ['Company Name', 'Market Capitalization', 'Last Sale']
print(listings.loc[tickers,info_cols].sort_values('Market Capitalization', ascending=False))
```

Sector	Stock Symbol	
Miscellaneous	BABA	275525.000000
Technology	ORCL	181046.096000
Health Care	ABBV	102196.076208
Transportation	UPS	90180.886756
Finance	GS	88840.590477
Consumer Non-Durables	ABEV	88240.198455
Basic Industries	RIO	70431.476895
Public Utilities	TEF	54609.806092
Capital Goods	GM	50086.335099
Consumer Services	LVS	44384.295569
Energy	PAA	22223.001416
Consumer Durables	WRK	12354.903312

Name: Market Capitalization, dtype: float64

```
Index(['RIO', 'GM', 'WRK', 'ABEV', 'LVS', 'PAA', 'GS', 'ABBV', 'BABA', 'TEF',
      'ORCL', 'UPS'],
      dtype='object', name='Stock Symbol')
```

Stock Symbol	Company Name	Market Capitalization \
BABA	Alibaba Group Holding Limited	275525.000000
BABA	Alibaba Group Holding Limited	275525.000000
BABA	Alibaba Group Holding Limited	275525.000000
ORCL	Oracle Corporation	181046.096000
ORCL	Oracle Corporation	181046.096000
ORCL	Oracle Corporation	181046.096000
ABBV	AbbVie Inc.	102196.076208
ABBV	AbbVie Inc.	102196.076208
ABBV	AbbVie Inc.	102196.076208
UPS	United Parcel Service, Inc.	90180.886756
UPS	United Parcel Service, Inc.	90180.886756
UPS	United Parcel Service, Inc.	90180.886756
GS	Goldman Sachs Group, Inc. (The)	88840.590477
GS	Goldman Sachs Group, Inc. (The)	88840.590477
GS	Goldman Sachs Group, Inc. (The)	88840.590477
ABEV	Ambev S.A.	88240.198455
ABEV	Ambev S.A.	88240.198455
ABEV	Ambev S.A.	88240.198455
RIO	Rio Tinto Plc	70431.476895
RIO	Rio Tinto Plc	70431.476895
RIO	Rio Tinto Plc	70431.476895
TEF	Telefonica SA	54609.806092
TEF	Telefonica SA	54609.806092
TEF	Telefonica SA	54609.806092
GM	General Motors Company	50086.335099
GM	General Motors Company	50086.335099
GM	General Motors Company	50086.335099
LVS	Las Vegas Sands Corp.	44384.295569
LVS	Las Vegas Sands Corp.	44384.295569
LVS	Las Vegas Sands Corp.	44384.295569
PAA	Plains All American Pipeline, L.P.	22223.001416

PAA	Plains All American Pipeline, L.P.	22223.001416
PAA	Plains All American Pipeline, L.P.	22223.001416
WRK	Westrock Company	12354.903312
WRK	Westrock Company	12354.903312
WRK	Westrock Company	12354.903312

Stock Symbol Last Sale

BABA	110.21
BABA	110.21
BABA	110.21
ORCL	44.00
ORCL	44.00
ORCL	44.00
ABBV	64.13
ABBV	64.13
ABBV	64.13
UPS	103.74
UPS	103.74
UPS	103.74
GS	223.32
GS	223.32
GS	223.32
ABEV	5.62
ABEV	5.62
ABEV	5.62
RIO	38.94
RIO	38.94
RIO	38.94
TEF	10.84
TEF	10.84
TEF	10.84
GM	33.39
GM	33.39
GM	33.39
LVS	55.90
LVS	55.90
LVS	55.90
PAA	30.72
PAA	30.72
PAA	30.72
WRK	49.34
WRK	49.34
WRK	49.34

Import index component price information

Now you'll use the stock symbols for the companies you selected in the last exercise to calculate returns for each company.

```
In [ ]: tickers = tickers.tolist()
```

```
In [ ]: print(tickers)
```

```
# Import prices and inspect result
```

```
stock_prices = pd.read_csv('C:\\Users\\yeiso\\OneDrive - Douglas College\\0. DOUGLAS COLLEGE\\3. Fund Machine Learning\\0. Python  
print(stock_prices.info())
```

```
# Calculate the returns
```

```
price_return = stock_prices.iloc[-1].div(stock_prices.iloc[0]).sub(1).mul(100)
```

```
# Plot horizontal bar chart of sorted price_return
```

```
price_return.sort_values().plot(kind='barh', title='Stock Price Returns');
```

```
['RIO', 'GM', 'WRK', 'ABEV', 'LVS', 'PAA', 'GS', 'ABBV', 'BABA', 'TEF', 'ORCL', 'UPS']
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
DatetimeIndex: 1762 entries, 2010-01-04 to 2016-12-30
```


```
Data columns (total 12 columns):
```

#	Column	Non-Null Count	Dtype
0	AAPL	1761 non-null	float64
1	AMGN	1761 non-null	float64
2	AMZN	1761 non-null	float64
3	CPRT	1761 non-null	float64
4	EL	1762 non-null	float64
5	GS	1762 non-null	float64
6	ILMN	1761 non-null	float64
7	MA	1762 non-null	float64
8	PAA	1762 non-null	float64
9	RIO	1762 non-null	float64
10	TEF	1762 non-null	float64
11	UPS	1762 non-null	float64

```
dtypes: float64(12)
```

```
memory usage: 179.0 KB
```

```
None
```



Build a market-cap weighted index

Calculate number of shares outstanding The next step towards building a value-weighted index is to calculate the number of shares for each index component.

The number of shares will allow you to calculate the total market capitalization for each component given the historical price series in the next exercise.

```
In [ ]: print(listings.info())
        print(tickers)

# Select components and relevant columns from listings
components = listings[['Market Capitalization', 'Last Sale']].loc[tickers]

# Print the first rows of components
print(components.head(5))

# Calculate the number of shares here
no_shares = components['Market Capitalization'].div(components['Last Sale'])
```

```
# Print the sorted no_shares  
print(no_shares.sort_values(ascending=False))
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
Index: 2901 entries, WBAI to ZTO
```

```
Data columns (total 7 columns):
```

#	Column	Non-Null Count	Dtype
0	Company Name	2901 non-null	object
1	Last Sale	2901 non-null	float64
2	Market Capitalization	2901 non-null	float64
3	IPO Year	2901 non-null	float64
4	Sector	2901 non-null	object
5	Industry	2901 non-null	object
6	Exchange	2901 non-null	object

```
dtypes: float64(3), object(4)
```

```
memory usage: 245.9+ KB
```

```
None
```

```
['RIO', 'GM', 'WRK', 'ABEV', 'LVS', 'PAA', 'GS', 'ABBV', 'BABA', 'TEF', 'ORCL', 'UPS']
```

```
Market Capitalization Last Sale
```

```
Stock Symbol
```

RIO	70431.476895	38.94
RIO	70431.476895	38.94
RIO	70431.476895	38.94
GM	50086.335099	33.39
GM	50086.335099	33.39

```
Stock Symbol
```

ABEV	15701.102928
ABEV	15701.102928
ABEV	15701.102928
TEF	5037.804990
TEF	5037.804990
TEF	5037.804990
ORCL	4114.684000
ORCL	4114.684000
ORCL	4114.684000
BABA	2500.000000
BABA	2500.000000
BABA	2500.000000
RIO	1808.717948
RIO	1808.717948
RIO	1808.717948
ABBV	1593.576738
ABBV	1593.576738
ABBV	1593.576738
GM	1500.039985
GM	1500.039985
GM	1500.039985
UPS	869.297154
UPS	869.297154
UPS	869.297154
LVS	793.994554
LVS	793.994554


```
LVS      793.994554
PAA      723.404994
PAA      723.404994
PAA      723.404994
GS       397.817439
GS       397.817439
GS       397.817439
WRK      250.403391
WRK      250.403391
WRK      250.403391
dtype: float64
```

Create time series of market value

You can now use the number of shares to calculate the total market capitalization for each component and trading date from the historical price series.

The result will be the key input to construct the value-weighted stock index, which you will complete in the next exercise.

```
In [ ]: components['Number of Shares'] = no_shares
        print(no_shares.sort_values())
```

Stock	Symbol
WRK	250.403391
WRK	250.403391
WRK	250.403391
GS	397.817439
GS	397.817439
GS	397.817439
PAA	723.404994
PAA	723.404994
PAA	723.404994
LVS	793.994554
LVS	793.994554
LVS	793.994554
UPS	869.297154
UPS	869.297154
UPS	869.297154
GM	1500.039985
GM	1500.039985
GM	1500.039985
ABBV	1593.576738
ABBV	1593.576738
ABBV	1593.576738
RIO	1808.717948
RIO	1808.717948
RIO	1808.717948
BABA	2500.000000
BABA	2500.000000
BABA	2500.000000
ORCL	4114.684000
ORCL	4114.684000
ORCL	4114.684000
TEF	5037.804990
TEF	5037.804990
TEF	5037.804990
ABEV	15701.102928
ABEV	15701.102928
ABEV	15701.102928

dtype: float64


```
In [ ]: no_shares_no_duplicates = no_shares.drop_duplicates()
print(no_shares_no_duplicates.sort_values())
```

```
Stock Symbol
WRK      250.403391
GS       397.817439
PAA      723.404994
LVS      793.994554
UPS      869.297154
GM       1500.039985
ABBV     1593.576738
RIO      1808.717948
BABA     2500.000000
ORCL     4114.684000
TEF      5037.804990
ABEV     15701.102928
dtype: float64
```

```
In [ ]: # Create the series of market cap per ticker
market_cap = stock_prices.mul(no_shares_no_duplicates)

# Select first and last market cap here
first_value = market_cap.iloc[0]
last_value = market_cap.iloc[-1]

# Concatenate and plot first and last market cap here
pd.concat([first_value, last_value], axis=1).dropna().plot(kind='barh');
plt.savefig('C:\\Users\\yeiso\\OneDrive - Douglas College\\0. DOUGLAS COLLEGE\\3. Fund Machine Learning\\0. Python Course DataCamp
plt.show()
```



Calculate & plot the composite index

By now you have all ingredients that you need to calculate the aggregate stock performance for your group of companies.

Use the time series of market capitalization that you created in the last exercise to aggregate the market value for each period, and then normalize this series to convert it to an index.

```
In [ ]: market_cap_series = market_cap[pd.concat([first_value, last_value], axis=1).dropna().index.tolist()]
```

```
In [ ]: market_cap_series
```

Out[]:

	GS	PAA	RIO	TEF	UPS
Date					
2010-01-04	68854.242342	19531.934838	101342.466626	143829.332464	50575.708420
2010-01-05	70071.563705	19748.956336	102916.051241	143728.576365	50662.638135
2010-01-06	69323.666920	19741.722286	106063.220471	142217.234868	50288.840359
2010-01-07	70680.224387	19502.998638	106081.307650	139799.088472	49906.349611
2010-01-08	69343.557792	19568.105088	107256.974316	138892.283574	52305.609756
...
2016-12-23	95862.068276	24168.960850	68948.328178	46196.671758	100812.390949
2016-12-27	96096.780565	24226.833249	69364.333306	45995.159559	100951.478494
2016-12-28	95734.766695	23908.535052	70304.866639	45491.379060	100143.032141
2016-12-29	94752.157621	23488.960155	70359.128177	45995.159559	99951.786767
2016-12-30	95257.385769	23358.747256	69563.292280	46347.805908	99656.225735


1762 rows × 5 columns

```
In [ ]: # Aggregate and print the market cap per trading day
raw_index = market_cap_series.sum(axis=1)
print(raw_index)

# Normalize the aggregate market cap here
index = raw_index.div(raw_index.iloc[0]).mul(100)
print(index)

# Plot the index here
index.plot(title='Market-Cap Weighted Index')
plt.show();
```

```
Date
2010-01-04    384133.684691
2010-01-05    387127.785783
2010-01-06    387634.684904
2010-01-07    385969.968759
2010-01-08    387366.530527
...
2016-12-23    335988.420011
2016-12-27    336634.585172
2016-12-28    335582.579586
2016-12-29    334547.192279
2016-12-30    334183.456947
Length: 1762, dtype: float64
Date
2010-01-04    100.000000
2010-01-05    100.779442
2010-01-06    100.911402
2010-01-07    100.478033
2010-01-08    100.841594
...
2016-12-23     87.466534
2016-12-27     87.634748
2016-12-28     87.360883
2016-12-29     87.091345
2016-12-30     86.996655
Length: 1762, dtype: float64
```



Evaluate index performance

- Index return:
 - Total index return
 - Contribution by component
- Performance vs Benchmark -Total period return
 - Rolling returns for sub periods

Calculate the contribution of each stock to the index

You have successfully built the value-weighted index. Let's now explore how it performed over the 2010-2016 period.

Let's also determine how much each stock has contributed to the index return.

```
In [ ]: # Calculate and print the index return here
index_return = (index.iloc[-1] / index.iloc[0] - 1) * 100
print(index_return)
```

```
# Select the market capitalization
market_cap = components['Market Capitalization']

# Calculate the total market cap
total_market_cap = market_cap.sum()

# Calculate the component weights , and print the result
weights = market_cap.div(total_market_cap)
print(weights.sort_values())

# Calculate and plot the distribution by component
weights.mul(index_return).sort_values().plot(kind='barh')
plt.show();
```


-13.003344859886202

Stock Symbol

WRK	0.003813
WRK	0.003813
WRK	0.003813
PAA	0.006858
PAA	0.006858
PAA	0.006858
LVS	0.013697
LVS	0.013697
LVS	0.013697
GM	0.015457
GM	0.015457
GM	0.015457
TEF	0.016853
TEF	0.016853
TEF	0.016853
RIO	0.021736
RIO	0.021736
RIO	0.021736
ABEV	0.027232
ABEV	0.027232
ABEV	0.027232
GS	0.027417
GS	0.027417
GS	0.027417
UPS	0.027831
UPS	0.027831
UPS	0.027831
ABBV	0.031539
ABBV	0.031539
ABBV	0.031539
ORCL	0.055872
ORCL	0.055872
ORCL	0.055872
BABA	0.085029
BABA	0.085029
BABA	0.085029

Name: Market Capitalization, dtype: float64


```
data.plot()  
plt.show();
```

```
Index    -13.003345  
DJIA      86.722172  
dtype: float64
```



Compare index performance against benchmark II

The next step in analyzing the performance of your index is to compare it against a benchmark.

In the video, we have use the S&P 500 as benchmark. You can also use the Dow Jones Industrial Average, which contains the 30 largest stocks, and would also be a reasonable benchmark for the largest stocks from all sectors across the three exchanges.

```
In [ ]: print(data.info())  
        print(data.head(5))  
  
# Create multi_period_return function here  
def multi_period_return(r):  
    return (np.prod(r + 1) - 1) * 100  
  
# Calculate rolling_return_360  
rolling_return_360 = data.pct_change().rolling('360D').apply(multi_period_return)
```


```
# Plot rolling_return_360 here
rolling_return_360.plot(title='Rolling 360D Return');
plt.show()
```

```
<class 'pandas.core.frame.DataFrame'>
DatetimeIndex: 42 entries, 2007-01-01 to 2017-04-01
Freq: QS-OCT
Data columns (total 2 columns):
#   Column  Non-Null Count  Dtype
---  -
0    gdp      41 non-null    float64
1   djia      40 non-null    float64
dtypes: float64(2)
memory usage: 1008.0 bytes
None
```

	gdp	djia
date		
2007-01-01	0.2	NaN
2007-04-01	3.1	NaN
2007-07-01	2.7	0.945735
2007-10-01	1.4	4.079072
2008-01-01	-2.7	-7.407889

C:\Users\yeiso\AppData\Local\Temp\ipykernel_19956\4027336615.py:9: FutureWarning: The default fill_method='pad' in DataFrame.pct_change is deprecated and will be removed in a future version. Call ffill before calling pct_change to retain current behavior and silence this warning.

```
rolling_return_360 = data.pct_change().rolling('360D').apply(multi_period_return)
```



Index correlation & exporting to Excel

Visualize your index constituent correlations To better understand the characteristics of your index constituents, you can calculate the return correlations.

Use the daily stock prices or your index companies, and show a heatmap of the daily return correlations!

```
In [ ]: stock_prices = pd.read_csv('C:\\Users\\yeiso\\OneDrive - Douglas College\\0. DOUGLAS COLLEGE\\3. Fund Machine Learning\\0. Python  
print(stock_prices.head())
```

	AAPL	AMGN	AMZN	CPRT	EL	GS	ILMN	MA	PAA	\
Date										
2010-01-04	30.57	57.72	133.90	4.55	24.27	173.08	30.55	25.68	27.00	
2010-01-05	30.63	57.22	134.69	4.55	24.18	176.14	30.35	25.61	27.30	
2010-01-06	30.14	56.79	132.25	4.53	24.25	174.26	32.22	25.56	27.29	
2010-01-07	30.08	56.27	130.00	4.50	24.56	177.67	32.77	25.39	26.96	
2010-01-08	30.28	56.77	133.52	4.52	24.66	174.31	33.15	25.40	27.05	

	RIO	TEF	UPS
Date			
2010-01-04	56.03	28.55	58.18
2010-01-05	56.90	28.53	58.28
2010-01-06	58.64	28.23	57.85
2010-01-07	58.65	27.75	57.41
2010-01-08	59.30	27.57	60.17

```
In [ ]: # Inspect stock_prices here
print(stock_prices.info())

# Calculate the daily returns
returns = stock_prices.pct_change()

# Calculate and print the pairwise correlations
correlations = returns.corr()
print(correlations)

# Plot a heatmap of daily return correlations
sns.heatmap(correlations, annot=True)
plt.title('Daily Return Correlations')
plt.show();
```

```

<class 'pandas.core.frame.DataFrame'>
DatetimeIndex: 1762 entries, 2010-01-04 to 2016-12-30
Data columns (total 12 columns):
#   Column  Non-Null Count  Dtype
---  -
0    AAPL      1761 non-null    float64
1    AMGN      1761 non-null    float64
2    AMZN      1761 non-null    float64
3    CPRT      1761 non-null    float64
4    EL        1762 non-null    float64
5    GS        1762 non-null    float64
6    ILMN      1761 non-null    float64
7    MA        1762 non-null    float64
8    PAA       1762 non-null    float64
9    RIO       1762 non-null    float64
10   TEF       1762 non-null    float64
11   UPS       1762 non-null    float64
dtypes: float64(12)
memory usage: 179.0 KB
None


```

	AAPL	AMGN	AMZN	CPRT	EL	GS	ILMN	\
AAPL	1.000000	0.286898	0.327611	0.346616	0.306770	0.344981	0.264791	
AMGN	0.286898	1.000000	0.323408	0.355892	0.349893	0.390076	0.336927	
AMZN	0.327611	0.323408	1.000000	0.298929	0.334031	0.333402	0.242726	
CPRT	0.346616	0.355892	0.298929	1.000000	0.371763	0.423160	0.265665	
EL	0.306770	0.349893	0.334031	0.371763	1.000000	0.358318	0.214027	
GS	0.344981	0.390076	0.333402	0.423160	0.358318	1.000000	0.266063	
ILMN	0.264791	0.336927	0.242726	0.265665	0.214027	0.266063	1.000000	
MA	0.391421	0.400230	0.428330	0.401352	0.431556	0.466796	0.301392	
PAA	0.212960	0.229255	0.182438	0.221273	0.206056	0.271982	0.162796	
RIO	0.361684	0.313878	0.326229	0.384944	0.415416	0.527298	0.234445	
TEF	0.325309	0.374555	0.331867	0.376767	0.428925	0.498230	0.231173	
UPS	0.366039	0.432468	0.378399	0.462716	0.456952	0.506407	0.267801	

	MA	PAA	RIO	TEF	UPS
AAPL	0.391421	0.212960	0.361684	0.325309	0.366039
AMGN	0.400230	0.229255	0.313878	0.374555	0.432468
AMZN	0.428330	0.182438	0.326229	0.331867	0.378399
CPRT	0.401352	0.221273	0.384944	0.376767	0.462716
EL	0.431556	0.206056	0.415416	0.428925	0.456952
GS	0.466796	0.271982	0.527298	0.498230	0.506407
ILMN	0.301392	0.162796	0.234445	0.231173	0.267801
MA	1.000000	0.243761	0.437778	0.448438	0.486512
PAA	0.243761	1.000000	0.337448	0.253598	0.217523
RIO	0.437778	0.337448	1.000000	0.559264	0.509809
TEF	0.448438	0.253598	0.559264	1.000000	0.516242
UPS	0.486512	0.217523	0.509809	0.516242	1.000000

C:\Users\yeiso\AppData\Local\Temp\ipykernel_19956\2553049458.py:5: FutureWarning: The default fill_method='pad' in DataFrame.pct_change is deprecated and will be removed in a future version. Call ffill before calling pct_change to retain current behavior and silence this warning.

```
returns = stock_prices.pct_change()
```



Save your analysis to multiple excel worksheets

Now that you have completed your analysis, you may want to save all results into a single Excel workbook.

Let's practice exporting various DataFrame to multiple Excel worksheets.

```
In [ ]: index = index.to_frame('Index')
```

```
In [ ]: print(index.info())
print(stock_prices.info())

# Join index to stock_prices, and inspect the result
data = stock_prices.join(index)
print(data.info())
```



```
# Create index & stock price returns
```

```
returns = data.pct_change()
```

```
#esta es la forma original.... pero NO funciona!
```

```
# Export data and data as returns to excel
```

```
#with pd.ExcelWriter('data.xls') as writer:
```

```
#    data.to_excel(writer, sheet_name='data')
```

```
#    returns.to_excel(writer, sheet_name='returns')
```

```
# Export data and data as returns to Excel with 'xls' format using 'openpyxl' engine
```

```
with pd.ExcelWriter('data.xls', engine='openpyxl') as writer:
```

```
    data.to_excel(writer, sheet_name='data', engine='openpyxl')
```

```
    returns.to_excel(writer, sheet_name='returns', engine='openpyxl')
```

```
<class 'pandas.core.frame.DataFrame'>
DatetimeIndex: 1762 entries, 2010-01-04 to 2016-12-30
Data columns (total 1 columns):
#   Column  Non-Null Count  Dtype
---  -
0    Index    1762 non-null    float64
dtypes: float64(1)
memory usage: 27.5 KB
None
```

```
<class 'pandas.core.frame.DataFrame'>
DatetimeIndex: 1762 entries, 2010-01-04 to 2016-12-30
Data columns (total 12 columns):
#   Column  Non-Null Count  Dtype
---  -
0    AAPL    1761 non-null    float64
1    AMGN    1761 non-null    float64
2    AMZN    1761 non-null    float64
3    CPRT    1761 non-null    float64
4    EL       1762 non-null    float64
5    GS       1762 non-null    float64
6    ILMN    1761 non-null    float64
7    MA       1762 non-null    float64
8    PAA      1762 non-null    float64
9    RIO      1762 non-null    float64
10   TEF      1762 non-null    float64
11   UPS      1762 non-null    float64
dtypes: float64(12)
memory usage: 179.0 KB
None
```

```
<class 'pandas.core.frame.DataFrame'>
DatetimeIndex: 1762 entries, 2010-01-04 to 2016-12-30
Data columns (total 13 columns):
#   Column  Non-Null Count  Dtype
---  -
0    AAPL    1761 non-null    float64
1    AMGN    1761 non-null    float64
2    AMZN    1761 non-null    float64
3    CPRT    1761 non-null    float64
4    EL       1762 non-null    float64
5    GS       1762 non-null    float64
6    ILMN    1761 non-null    float64
7    MA       1762 non-null    float64
8    PAA      1762 non-null    float64
9    RIO      1762 non-null    float64
10   TEF      1762 non-null    float64
11   UPS      1762 non-null    float64
12   Index    1762 non-null    float64
dtypes: float64(13)
memory usage: 257.3 KB
None
```

```
C:\Users\yeiso\AppData\Local\Temp\ipykernel_19956\1964733718.py:9: FutureWarning: The default fill_method='pad' in DataFrame.pct_change is deprecated and will be removed in a future version. Call ffill before calling pct_change to retain current behavior and silence this warning.
```

```
    returns = data.pct_change()
```