

Introduction to preprocessing

PREPROCESSING FOR MACHINE LEARNING IN PYTHON



James Chapman
Curriculum Manager, DataCamp

What is data preprocessing?

- After exploratory data analysis and data cleaning
- Preparing data for modeling
- **Example:** transforming categorical features into numerical features (dummy variables)

Why preprocess?

- Transform dataset so it's suitable for modeling
- Improve model performance
- Generate more reliable results



Recap: exploring data with pandas

```
import pandas as pd
hiking = pd.read_json("hiking.json")
print(hiking.head())
```

```
   Prop_ID      Name  ... lat lon
0    B057  Salt Marsh Nature Trail  ... NaN NaN
1    B073      Lullwater  ... NaN NaN
2    B073      Midwood  ... NaN NaN
3    B073    Peninsula  ... NaN NaN
4    B073    Waterfall  ... NaN NaN
```

Recap: exploring data with pandas

```
print(hiking.info())
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 33 entries, 0 to 32
Data columns (total 11 columns):
#   Column          Non-Null Count  Dtype
--  --
0   Prop_ID         33 non-null    object
1   Name            33 non-null    object
2   Location        33 non-null    object
3   Park_Name       33 non-null    object
4   Length          29 non-null    object
5   Difficulty       27 non-null    object
6   Other_Details   31 non-null    object
7   Accessible      33 non-null    object
8   Limited_Access  33 non-null    object
9   lat             0 non-null     float64
10  lon             0 non-null     float64
dtypes: float64(2), object(9)
memory usage: 3.0+ KB
```

Recap: exploring data with pandas

```
print(wine.describe())
```

	Type	Alcohol	...	Alcalinity of ash
count	178.000000	178.000000	...	178.000000
mean	1.938202	13.000618	...	19.494944
std	0.775035	0.811827	...	3.339564
min	1.000000	11.030000	...	10.600000
25%	1.000000	12.362500	...	17.200000
50%	2.000000	13.050000	...	19.500000
75%	3.000000	13.677500	...	21.500000
max	3.000000	14.830000	...	30.000000

Removing missing data

```
print(df)
```

	A	B	C
0	1.0	NaN	2.0
1	4.0	7.0	3.0
2	7.0	NaN	NaN
3	NaN	7.0	NaN
4	5.0	9.0	7.0

```
print(df.dropna())
```

	A	B	C
1	4.0	7.0	3.0
4	5.0	9.0	7.0

Removing missing data

```
print(df)
```

	A	B	C
0	1.0	NaN	2.0
1	4.0	7.0	3.0
2	7.0	NaN	NaN
3	NaN	7.0	NaN
4	5.0	9.0	7.0

```
print(df.drop([1, 2, 3]))
```

	A	B	C
0	1.0	NaN	2.0
4	5.0	9.0	7.0

Removing missing data

```
print(df)
```

	A	B	C
0	1.0	NaN	2.0
1	4.0	7.0	3.0
2	7.0	NaN	NaN
3	NaN	7.0	NaN
4	5.0	9.0	7.0

```
print(df.drop("A", axis=1))
```

	B	C
0	NaN	2.0
1	7.0	3.0
2	NaN	NaN
3	7.0	NaN
4	9.0	7.0

Removing missing data

```
print(df)
```

	A	B	C
0	1.0	NaN	2.0
1	4.0	7.0	3.0
2	7.0	NaN	NaN
3	NaN	7.0	NaN
4	5.0	9.0	7.0

```
print(df.isna().sum())
```

A	1
B	2
C	2

dtype: int64

```
print(df.dropna(subset=["B"]))
```

	A	B	C
1	4.0	7.0	3.0
3	NaN	7.0	NaN
4	5.0	9.0	7.0

Removing missing data

```
print(df)
```

	A	B	C
0	1.0	NaN	2.0
1	4.0	7.0	3.0
2	7.0	NaN	NaN
3	NaN	7.0	NaN
4	5.0	9.0	7.0

```
print(df.dropna(thresh=2))
```

	A	B	C
0	1.0	NaN	2.0
1	4.0	7.0	3.0
4	5.0	9.0	7.0

Let's practice!

PREPROCESSING FOR MACHINE LEARNING IN PYTHON

Working With Data Types

PREPROCESSING FOR MACHINE LEARNING IN PYTHON



James Chapman
Curriculum Manager, DataCamp

Why are types important?

```
print(volunteer.info())
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 665 entries, 0 to 664
Data columns (total 35 columns):
#   Column                Non-Null Count  Dtype
--  --
0   opportunity_id        665 non-null   int64
1   content_id            665 non-null   int64
2   vol_requests          665 non-null   int64
3   event_time            665 non-null   int64
4   title                 665 non-null   object
..  ...
34  NTA                   0 non-null     float64
dtypes: float64(13), int64(8), object(14)
memory usage: 182.0+ KB
```

- `object` : string/mixed types
- `int64` : integer
- `float64` : float
- `datetime64` : dates and times

Converting column types

```
print(df)
```

	A	B	C
0	1	string	1.0
1	2	string2	2.0
2	3	string3	3.0

```
print(df.info())
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 3 entries, 0 to 2
Data columns (total 3 columns):
#   Column  Non-Null Count  Dtype
--  --
0    A      3 non-null      int64
1    B      3 non-null      object
2    C      3 non-null      object
dtypes: int64(1), object(2)
memory usage: 200.0+ bytes
```

Converting column types

```
print(df)
```

	A	B	C
0	1	string	1.0
1	2	string2	2.0
2	3	string3	3.0

```
df["C"] = df["C"].astype("float")  
print(df.dtypes)
```

```
A      int64  
B      object  
C      float64  
dtype: object
```


Let's practice!

PREPROCESSING FOR MACHINE LEARNING IN PYTHON

Training and test sets

PREPROCESSING FOR MACHINE LEARNING IN PYTHON



James Chapman
Curriculum Manager, DataCamp

Why split?

1. Reduces *overfitting*
2. Evaluate performance on a holdout set

Splitting up your dataset

```
from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
```

	X_train	y_train
0	1.0	n
1	4.0	n
	...	
5	5.0	n
6	6.0	n

	X_test	y_test
0	9.0	y
1	1.0	n
2	4.0	n

Stratified sampling

- Dataset of 100 samples: 80 **class 1** and 20 **class 2**
- Training set of 75 samples: 60 **class 1** and 15 **class 2**
- Test set of 25 samples: 20 **class 1** and 5 **class 2**

Stratified sampling

```
X_train, X_test, y_train, y_test = train_test_split(X, y, stratify=y, random_state=42)
```

```
y["labels"].value_counts()
```

```
class1    80  
class2    20  
Name: labels, dtype: int64
```

Stratified sampling

```
y_train["labels"].value_counts()
```

```
class1    60  
class2    15  
Name: labels, dtype: int64
```

```
y_test["labels"].value_counts()
```

```
class1     20  
class2      5  
Name: labels, dtype: int64
```

Let's practice!

PREPROCESSING FOR MACHINE LEARNING IN PYTHON

Standardization

PREPROCESSING FOR MACHINE LEARNING IN PYTHON



James Chapman

Curriculum Manager, DataCamp

What is standardization?

Standardization: transform *continuous* data to appear *normally distributed*


- `scikit-learn` models assume normally distributed data
- Using non-normal training data can introduce *bias*
- **Log normalization** and feature **scaling** in this course
- Applied to continuous numerical data

When to standardize: linear distances

- Model in *linear* space

Examples:

- k-Nearest Neighbors (kNN)
- Linear regression
- K-Means Clustering




When to standardize: high variance

- Model in *linear* space

Examples:

- k-Nearest Neighbors (kNN)
- Linear regression
- K-Means Clustering
- Dataset features have *high variance*



When to standardize: different scales

- Features are on *different scales*

Example:

- Predicting house prices using *no. bedrooms* and *last sale price*
- Linearity assumptions

Let's practice!

PREPROCESSING FOR MACHINE LEARNING IN PYTHON

Log normalization

PREPROCESSING FOR MACHINE LEARNING IN PYTHON



James Chapman

Curriculum Manager, DataCamp

What is log normalization?

- Useful for features with *high variance*
- Applies logarithm transformation
- Natural log using the constant e (≈ 2.718)

What is log normalization?

- Useful for features with *high variance*
 - Applies logarithm transformation
 - Natural log using the constant e (≈ 2.718)
 - $e^{3.4} = 30$
-
- Captures relative changes, the magnitude of change, and keeps everything positive

Number	Log
30	3.4
300	5.7
3000	8

Log normalization in Python

```
print(df)
```

```
   col1  col2
0  1.00   3.0
1  1.20  45.5
2  0.75  28.0
3  1.60 100.0
```

```
print(df.var())
```

```
col1      0.128958
col2    1691.729167
dtype: float64
```

```
import numpy as np
df["log_2"] = np.log(df["col2"])
print(df)
```

```
   col1  col2  log_2
0  1.00   3.0  1.098612
1  1.20  45.5  3.817712
2  0.75  28.0  3.332205
3  1.60 100.0  4.605170
```

```
print(df[["col1", "log_2"]].var())
```

```
col1      0.128958
log_2     2.262886
dtype: float64
```

Let's practice!

PREPROCESSING FOR MACHINE LEARNING IN PYTHON

Scaling data

PREPROCESSING FOR MACHINE LEARNING IN PYTHON



James Chapman

Curriculum Manager, DataCamp

What is feature scaling?

- Features on different scales
- Model with linear characteristics
- Center features around 0 and transform to variance of 1
- Transforms to approximately normal distribution

How to scale data

```
print(df)
```

```
   col1  col2  col3
0  1.00  48.0  100.0
1  1.20  45.5  101.3
2  0.75  46.2  103.5
3  1.60  50.0  104.0
```

```
print(df.var())
```

```
col1    0.128958
col2    4.055833
col3    3.526667
dtype: float64
```

How to scale data

```
from sklearn.preprocessing import StandardScaler
scaler = StandardScaler()
df_scaled = pd.DataFrame(scaler.fit_transform(df),
                          columns=df.columns)
```

```
print(df_scaled)
```

```
      col1      col2      col3
0 -0.442127  0.329683 -1.352726
1  0.200967 -1.103723 -0.553388
2 -1.245995 -0.702369  0.799338
3  1.487156  1.476409  1.106776
```

```
print(df_scaled.var())
```

```
col1    1.333333
col2    1.333333
col3    1.333333
dtype: float64
```

Let's practice!

PREPROCESSING FOR MACHINE LEARNING IN PYTHON

Standardized data and modeling

PREPROCESSING FOR MACHINE LEARNING IN PYTHON



James Chapman
Curriculum Manager, DataCamp

K-nearest neighbors

- **Data leakage:** non-training data is used to train the model

```
from sklearn.model_selection import train_test_split
from sklearn.neighbors import KNeighborsClassifier

X_train, X_test, y_train, y_test = train_test_split(X, y, stratify=y, random_state=42)
knn = KNeighborsClassifier()
scaler = StandardScaler()

X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)

knn.fit(X_train_scaled, y_train)
knn.score(X_test_scaled, y_test)
```

Let's practice!

PREPROCESSING FOR MACHINE LEARNING IN PYTHON

Feature engineering

PREPROCESSING FOR MACHINE LEARNING IN PYTHON



James Chapman

Curriculum Manager, DataCamp

What is feature engineering?

Feature engineering: Creation of new features from existing ones

- Improve performance
- Insight into relationships between features
- Need to understand the data first!
- Highly dataset-dependent

Feature engineering scenarios

Id	Text
1	"Feature engineering is fun!"
2	"Feature engineering is a lot of work."
3	"I don't mind feature engineering."

user	fav_color
1	blue
2	green
3	orange

Feature engineering scenarios

Id	Date
4	July 30 2011
5	January 29 2011
6	February 05 2011

user	test1	test2	test3
1	90.5	89.6	91.4
2	65.5	70.6	67.3
3	78.1	80.7	81.8

Let's practice!

PREPROCESSING FOR MACHINE LEARNING IN PYTHON

Encoding categorical variables

PREPROCESSING FOR MACHINE LEARNING IN PYTHON



James Chapman
Curriculum Manager, DataCamp

Categorical variables

```
   user subscribed fav_color
0     1         y    blue
1     2         n   green
2     3         n  orange
3     4         y   green
```

Encoding binary variables - pandas

```
print(users["subscribed"])
```

```
0    y
1    n
2    n
3    y
Name: subscribed, dtype: object
```

```
print(users[["subscribed", "sub_enc"]])
```

	subscribed	sub_enc
0	y	1
1	n	0
2	n	0
3	y	1

```
users["sub_enc"] = users["subscribed"].apply(lambda val: 1 if val == "y" else 0)
```

Encoding binary variables - scikit-learn

```
from sklearn.preprocessing import LabelEncoder

le = LabelEncoder()
users["sub_enc_le"] = le.fit_transform(users["subscribed"])

print(users[["subscribed", "sub_enc_le"]])
```

	subscribed	sub_enc_le
0	y	1
1	n	0
2	n	0
3	y	1

One-hot encoding

fav_color
blue
green
orange
green

Values: [blue, green, orange]

- blue: [1, 0, 0]
- green: [0, 1, 0]
- orange: [0, 0, 1]

fav_color_enc
[1, 0, 0]
[0, 1, 0]
[0, 0, 1]
[0, 1, 0]

```
print(users["fav_color"])
```

```
0    blue
1   green
2  orange
3   green
Name: fav_color, dtype: object
```

```
print(pd.get_dummies(users["fav_color"]))
```

	blue	green	orange
0	1	0	0
1	0	1	0
2	0	0	1
3	0	1	0

Let's practice!

PREPROCESSING FOR MACHINE LEARNING IN PYTHON

Engineering numerical features

PREPROCESSING FOR MACHINE LEARNING IN PYTHON



James Chapman
Curriculum Manager, DataCamp


```
print(temps)
```

```
   city  day1  day2  day3
0   NYC  68.3  67.9  67.8
1    SF  75.1  75.5  74.9
2    LA  80.3  84.0  81.3
3 Boston  63.0  61.0  61.2
```

```
temps["mean"] = temps.loc[:, "day1": "day3"].mean(axis=1)
print(temps)
```

```
   city  day1  day2  day3  mean
0   NYC  68.3  67.9  67.8  68.00
1    SF  75.1  75.5  74.9  75.17
2    LA  80.3  84.0  81.3  81.87
3 Boston  63.0  61.0  61.2  61.73
```

Dates

```
print(purchases)
```

```
      date purchase
0  July 30 2011  $45.08
1 February 01 2011  $19.48
2 January 29 2011  $76.09
3  March 31 2012  $32.61
4 February 05 2011  $75.98
```

Dates

```
purchases["date_converted"] = pd.to_datetime(purchases["date"])
purchases['month'] = purchases["date_converted"].dt.month
print(purchases)
```


			date	purchase	date_converted	month
0	July	30	2011	\$45.08	2011-07-30	7
1	February	01	2011	\$19.48	2011-02-01	2
2	January	29	2011	\$76.09	2011-01-29	1
3	March	31	2012	\$32.61	2012-03-31	3
4	February	05	2011	\$75.98	2011-02-05	2

Let's practice!

PREPROCESSING FOR MACHINE LEARNING IN PYTHON

Engineering text features

PREPROCESSING FOR MACHINE LEARNING IN PYTHON



James Chapman
Curriculum Manager, DataCamp

Extraction

- Regular expressions: code to identify patterns

```
import re
my_string = "temperature:75.6 F"
temp = re.search("\d+\.\d+", my_string)

print(float(temp.group(0)))
```

75.6

- \d+
- \.
- \d+

Vectorizing text

TF/IDF: Vectorizes words based upon importance

- TF = Term Frequency
- IDF = Inverse Document Frequency

Vectorizing text

```
from sklearn.feature_extraction.text import TfidfVectorizer
```

```
print(documents.head())
```

```
0    Building on successful events last summer and ...  
1           Build a website for an Afghan business  
2    Please join us and the students from Mott Hall...  
3    The Oxfam Action Corps is a group of dedicated...  
4    Stop 'N' Swap reduces NYC's waste by finding n...
```

```
tfidf_vec = TfidfVectorizer()  
text_tfidf = tfidf_vec.fit_transform(documents)
```


Text classification

$$P(A|B) = \frac{P(B|A)P(A)}{P(B)}$$

Let's practice!

PREPROCESSING FOR MACHINE LEARNING IN PYTHON

Feature selection

PREPROCESSING FOR MACHINE LEARNING IN PYTHON



James Chapman

Curriculum Manager, DataCamp

What is feature selection?

- Selecting features to be used for modeling
- Doesn't create new features
- Improve model's performance

When to select features

city	state	lat	long
hico	tx	31.982778	-98.033333
mackinaw city	mi	45.783889	-84.727778
winchester	ky	37.990000	-84.179722


- Reducing noise
- Features are strongly statistically correlated
- Reduce overall variance

Let's practice!

PREPROCESSING FOR MACHINE LEARNING IN PYTHON

Removing redundant features

PREPROCESSING FOR MACHINE LEARNING IN PYTHON



James Chapman
Curriculum Manager, DataCamp

Redundant features

- Remove noisy features
- Remove correlated features
- Remove duplicated features

Scenarios for manual removal

city	state	lat	long
hico	tx	31.982778	-98.033333
mackinaw city	mi	45.783889	-84.727778
winchester	ky	37.990000	-84.179722

Correlated features

- Statistically correlated: features move together directionally
- Linear models assume feature independence
- Pearson's correlation coefficient

Correlated features

```
print(df)
```

```
      A      B      C
0  3.06  3.92  1.04
1  2.76  3.40  1.05
2  3.24  3.17  1.03
...
```

```
print(df.corr())
```

```
      A      B      C
A  1.000000  0.787194  0.543479
B  0.787194  1.000000  0.565468
C  0.543479  0.565468  1.000000
```

Let's practice!

PREPROCESSING FOR MACHINE LEARNING IN PYTHON

Selecting features using text vectors

PREPROCESSING FOR MACHINE LEARNING IN PYTHON



James Chapman
Curriculum Manager, DataCamp

Looking at word weights

```
print(tfidf_vec.vocabulary_)
```

```
{'200': 0,  
 '204th': 1,  
 '33rd': 2,  
 'ahead': 3,  
 'alley': 4,  
 ...}
```

```
print(text_tfidf[3].data)
```

```
[0.19392702 0.20261085 ...]
```

```
print(text_tfidf[3].indices)
```

```
[ 31 102  20  70   5 ...]
```

Looking at word weights

```
vocab = {v:k for k,v in  
tfidf_vec.vocabulary_.items()}
```

```
print(vocab)
```

```
{0: '200',  
1: '204th',  
2: '33rd',  
3: 'ahead',  
4: 'alley',  
...
```

```
zipped_row = dict(zip(text_tfidf[3].indices,  
text_tfidf[3].data))
```

```
print(zipped_row)
```

```
{5: 0.1597882543332701,  
7: 0.26576432098763175,  
8: 0.18599931331925676,  
9: 0.26576432098763175,  
10: 0.13077355258450366,  
...
```

Looking at word weights

```
def return_weights(vocab, vector, vector_index):  
  
    zipped = dict(zip(vector[vector_index].indices,  
                      vector[vector_index].data))  
  
    return {vocab[i]:zipped[i] for i in vector[vector_index].indices}  
  
print(return_weights(vocab, text_tfidf, 3))
```

```
{'and': 0.1597882543332701,  
 'are': 0.26576432098763175,  
 'at': 0.18599931331925676,  
 ...}
```


Let's practice!

PREPROCESSING FOR MACHINE LEARNING IN PYTHON

Dimensionality reduction

PREPROCESSING FOR MACHINE LEARNING IN PYTHON



James Chapman
Curriculum Manager, DataCamp

Dimensionality reduction and PCA

- Unsupervised learning method
- Combines/decomposes a feature space
- Feature extraction - here we'll use to reduce our feature space
- Principal component analysis
- Linear transformation to uncorrelated space
- Captures as much variance as possible in each component

PCA in scikit-learn

```
from sklearn.decomposition import PCA
pca = PCA()
df_pca = pca.fit_transform(df)

print(df_pca)
```

```
[88.4583, 18.7764, -2.2379, ..., 0.0954, 0.0361, -0.0034],
[93.4564, 18.6709, -1.7887, ..., -0.0509, 0.1331, 0.0119],
[-186.9433, -0.2133, -5.6307, ..., 0.0332, 0.0271, 0.0055]
```

```
print(pca.explained_variance_ratio_)
```

```
[0.9981, 0.0017, 0.0001, 0.0001, ...]
```

PCA caveats

- Difficult to interpret components
- End of preprocessing journey

Let's practice!

PREPROCESSING FOR MACHINE LEARNING IN PYTHON

UFOs and preprocessing

PREPROCESSING FOR MACHINE LEARNING IN PYTHON



James Chapman
Curriculum Manager, DataCamp

Identifying areas for preprocessing



Important concepts to remember


- Missing data: `.dropna()` and `.isna()`
- Types: `.astype()`
- Stratified sampling: `train_test_split(X, y, stratify=y)`

Let's practice!

PREPROCESSING FOR MACHINE LEARNING IN PYTHON

Categorical variables and standardization

PREPROCESSING FOR MACHINE LEARNING IN PYTHON



James Chapman
Curriculum Manager, DataCamp

Categorical variables

```
state country      type
295    az      us    light
296    tx      us  formation
297    nv      us  fireball
```

- One-hot encoding: `pd.get_dummies()`

Standardization


- `.var()`
- `np.log()`

Let's practice!

PREPROCESSING FOR MACHINE LEARNING IN PYTHON

Engineering new features

PREPROCESSING FOR MACHINE LEARNING IN PYTHON



James Chapman
Curriculum Manager, DataCamp

UFO feature engineering

date	length_of_time	desc
6/16/2013 21:00	5 minutes	Sabino Canyon Tucson Arizona night UFO sighting.
9/12/2005 22:35	5 minutes	Star like objects hovering in sky, slowly m...
12/31/2013 22:25	3 minutes	Three orange fireballs spotted by witness in E...

- Dates: `.dt.month` or `.dt.hour` attributes
- Regex: `\d` and `.group()`
- Text: tf-idf and `TfidfVectorizer`

Let's practice!

PREPROCESSING FOR MACHINE LEARNING IN PYTHON

Feature selection and modeling

PREPROCESSING FOR MACHINE LEARNING IN PYTHON



James Chapman
Curriculum Manager, DataCamp

Feature selection and modeling

- Redundant features
- Text vector

Final thoughts

- Iterative processes
- Know your dataset
- Understand your modeling task

Let's practice!

PREPROCESSING FOR MACHINE LEARNING IN PYTHON

Congratulations!

PREPROCESSING FOR MACHINE LEARNING IN PYTHON



James Chapman

Curriculum Manager, DataCamp

What you've learned

- Preparing data for modeling:
 - Missing data
 - Incorrect types
 - Standardize numerical values
 - Process categorical values
 - Feature engineering
 - Select features for modeling

Let's practice!

PREPROCESSING FOR MACHINE LEARNING IN PYTHON